

Data Structures

Lecture 4

AVL and WAVL Trees

Haim Kaplan and Uri Zwick

November 2014

Last updated: April 16, 2018

Balanced search trees

$O(\log n)$ worst-case time for all operations

AVL trees (1962)

Red-Black trees (1972)

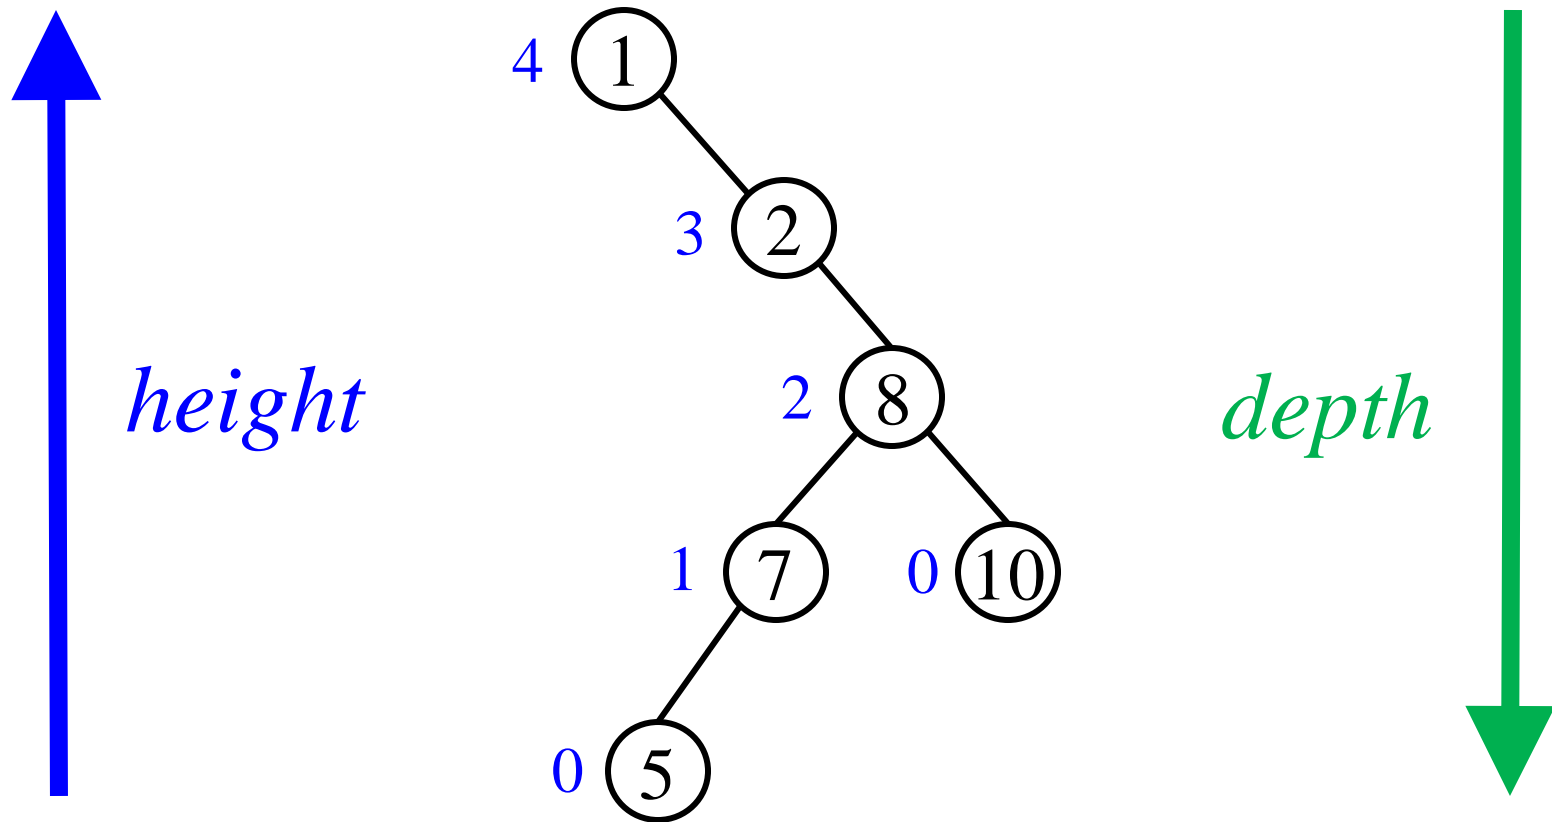
⋮

WAVL trees (2009)

Splay trees (amortized bounds)

B-trees (non-binary)

Height – Length of longest path to leaf



$$height(x) = 1 + \max\{height(x.left), height(x.right)\}$$

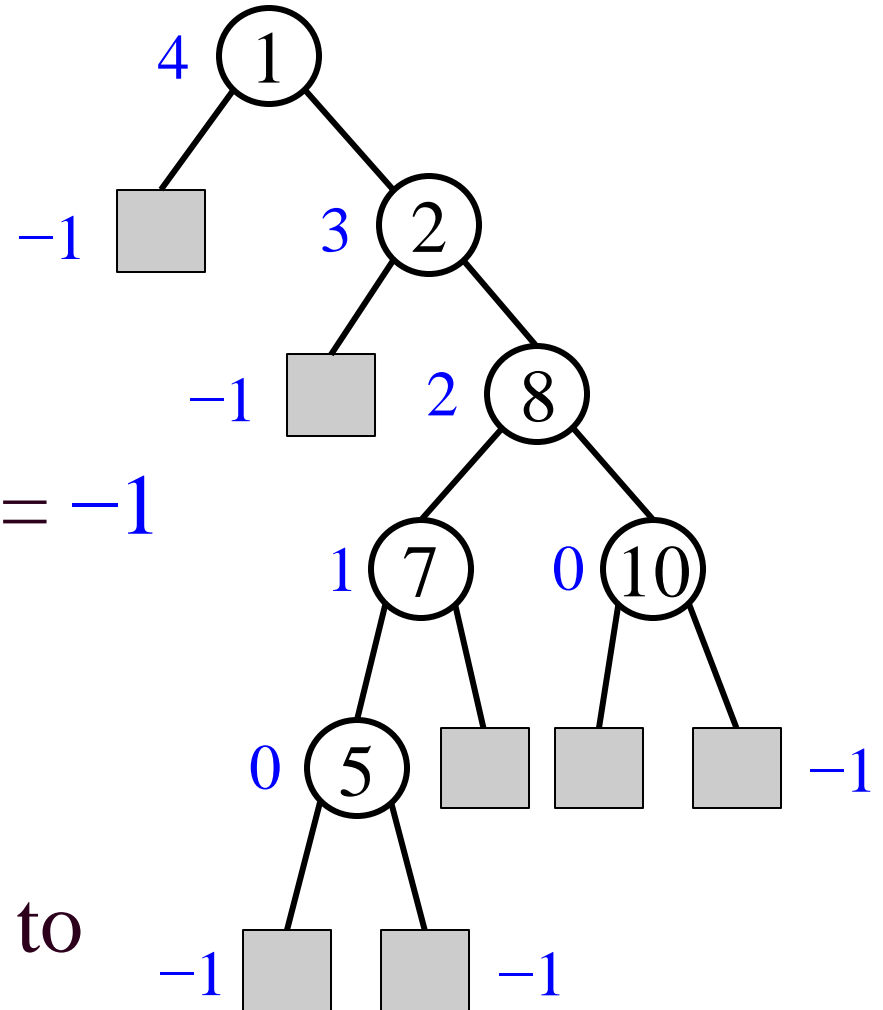
External leaves

Height of a leaf = 0

Height of a external leaf = -1

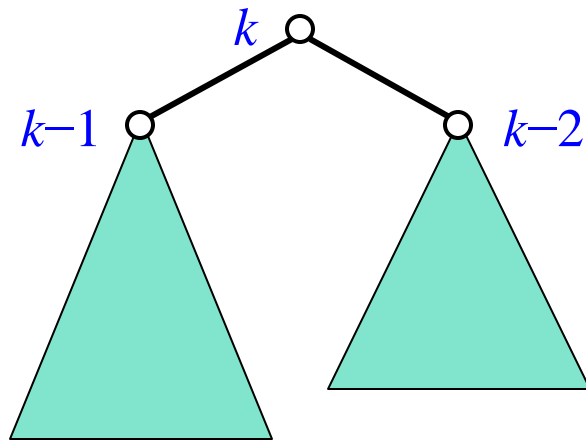
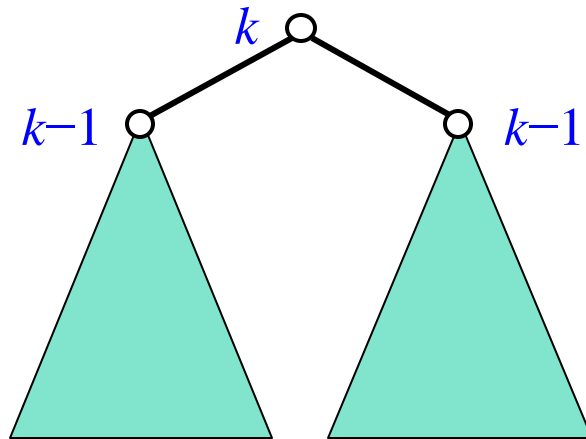
Unifies the treatment
of base cases

A single object EXT used to
represent all external leaves



AVL trees

[Adel'son-Vel'skii, Landis (1962)]



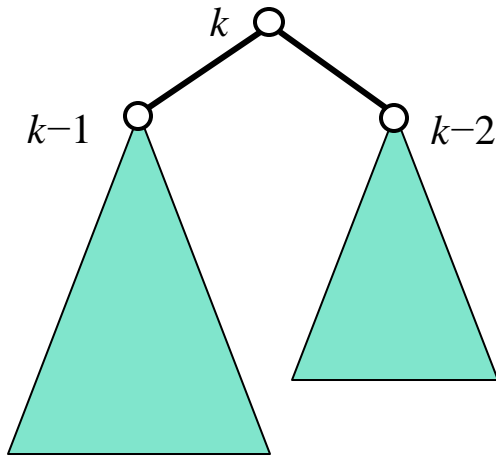
AVL trees are
search trees

The *height* of two
siblings differs by
at most 1

Need only one
extra bit per node

AVL trees

[Adel'son-Vel'skii, Landis (1962)]



S_k – minimal number of nodes
in an AVL tree of height k

$$S_k = S_{k-1} + S_{k-2} + 1$$

$$S_{-1} = 0, S_0 = 1$$

By induction: $S_k = F_{k+3} - 1 \geq F_{k+2} \geq \phi^k$

$$\text{height} \leq \log_{\phi} n \leq 1.4404 \log_2 n$$

Fibonacci numbers

$$F_k = F_{k-1} + F_{k-2} , \ k \geq 2$$

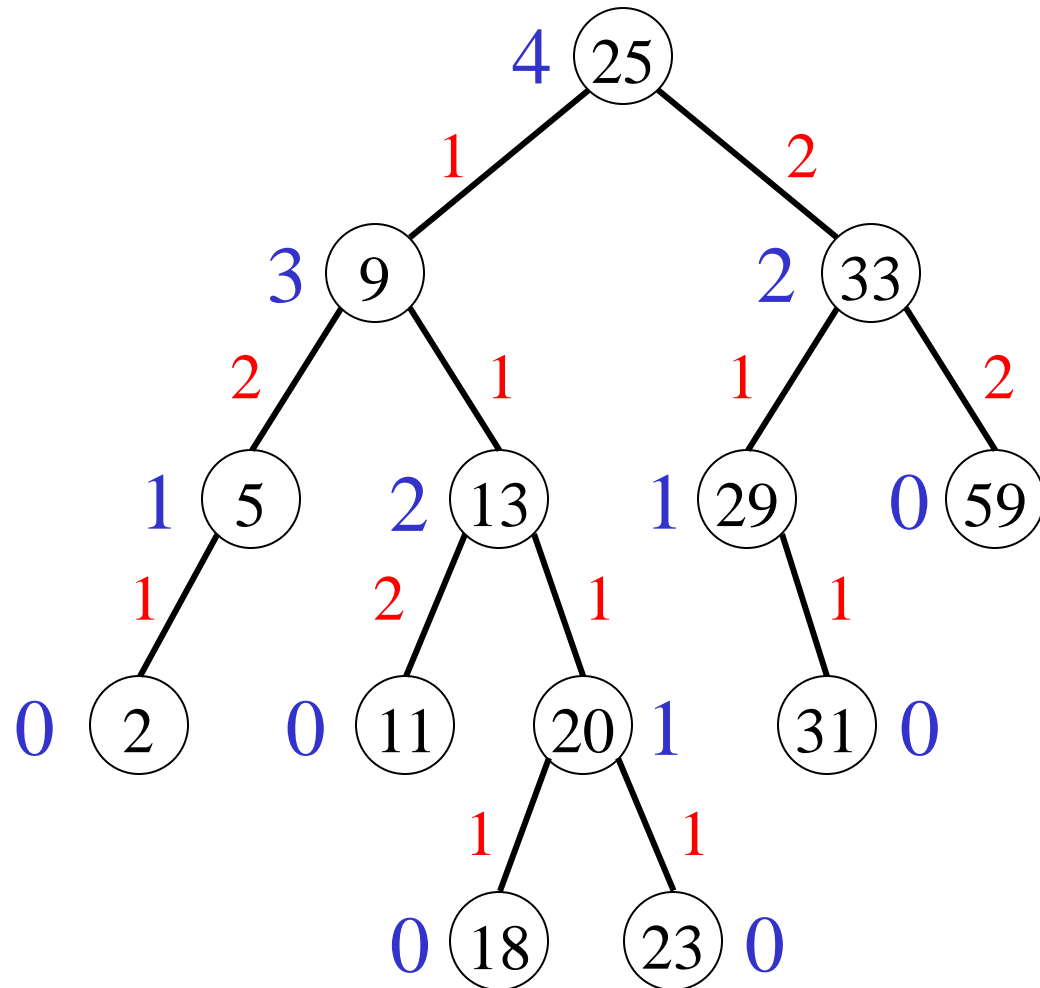
$$F_0 = 0 , \ F_1 = 1$$

n	0	1	2	3	4	5	6	7	8	9
F_n	0	1	1	2	3	5	8	13	21	34

$$F_k \geq \phi^{k-2}$$

$$\phi = \frac{1+\sqrt{5}}{2} \simeq 1.618$$

An AVL tree



The challenge

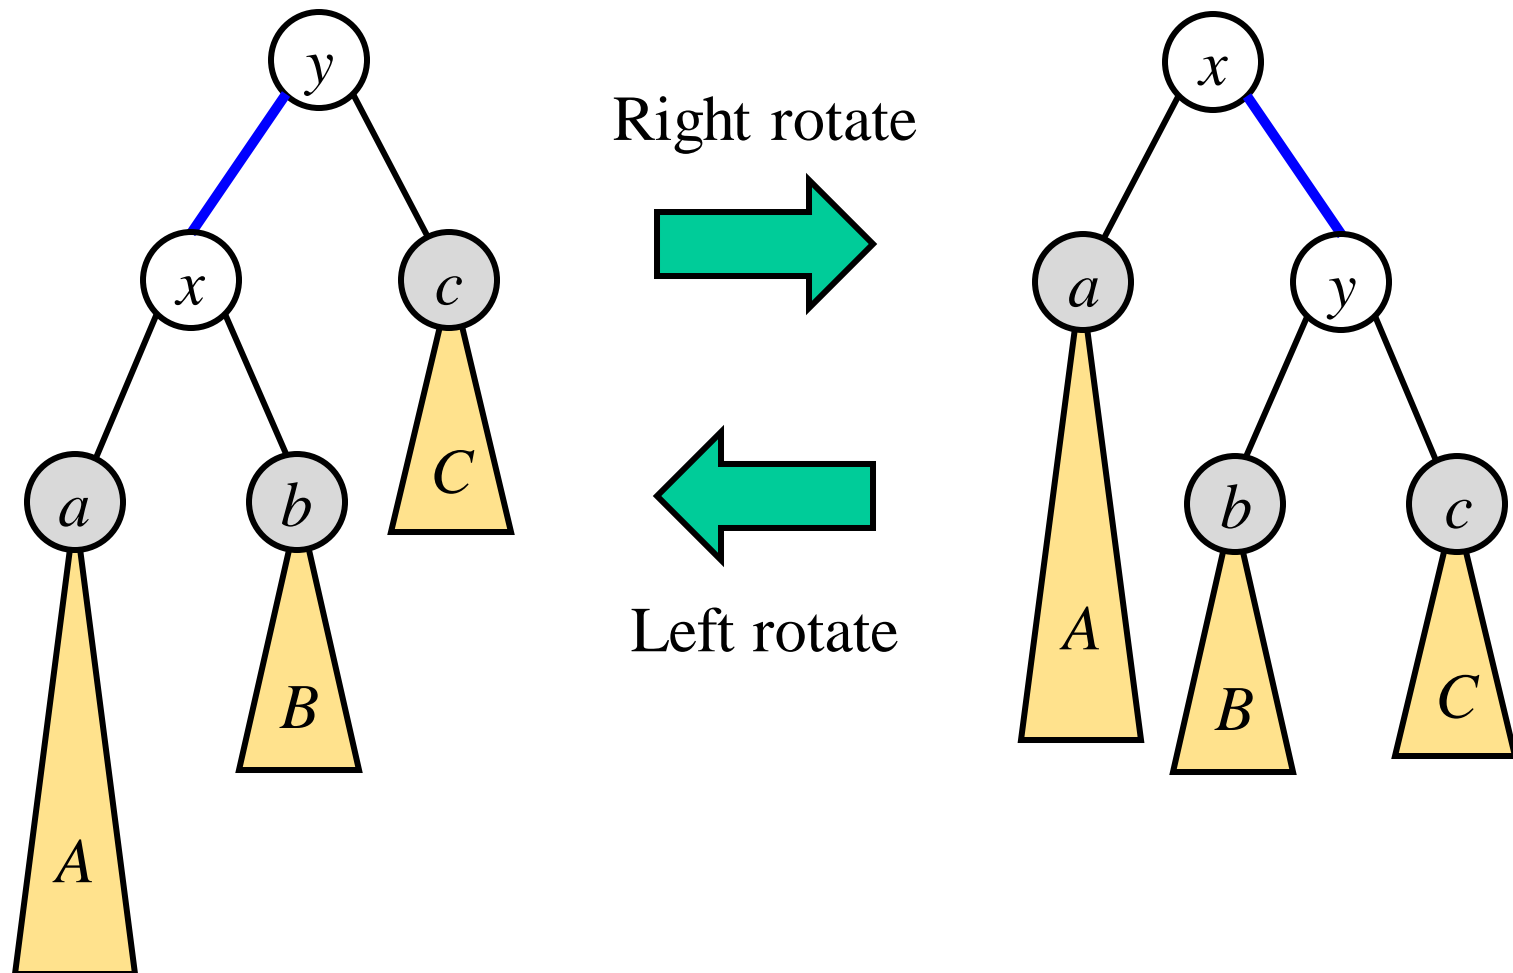
If we insert or delete a node from an AVL tree, the resulting tree is not necessarily an AVL tree

After insertions and deletions
we may need to **restructure** the tree

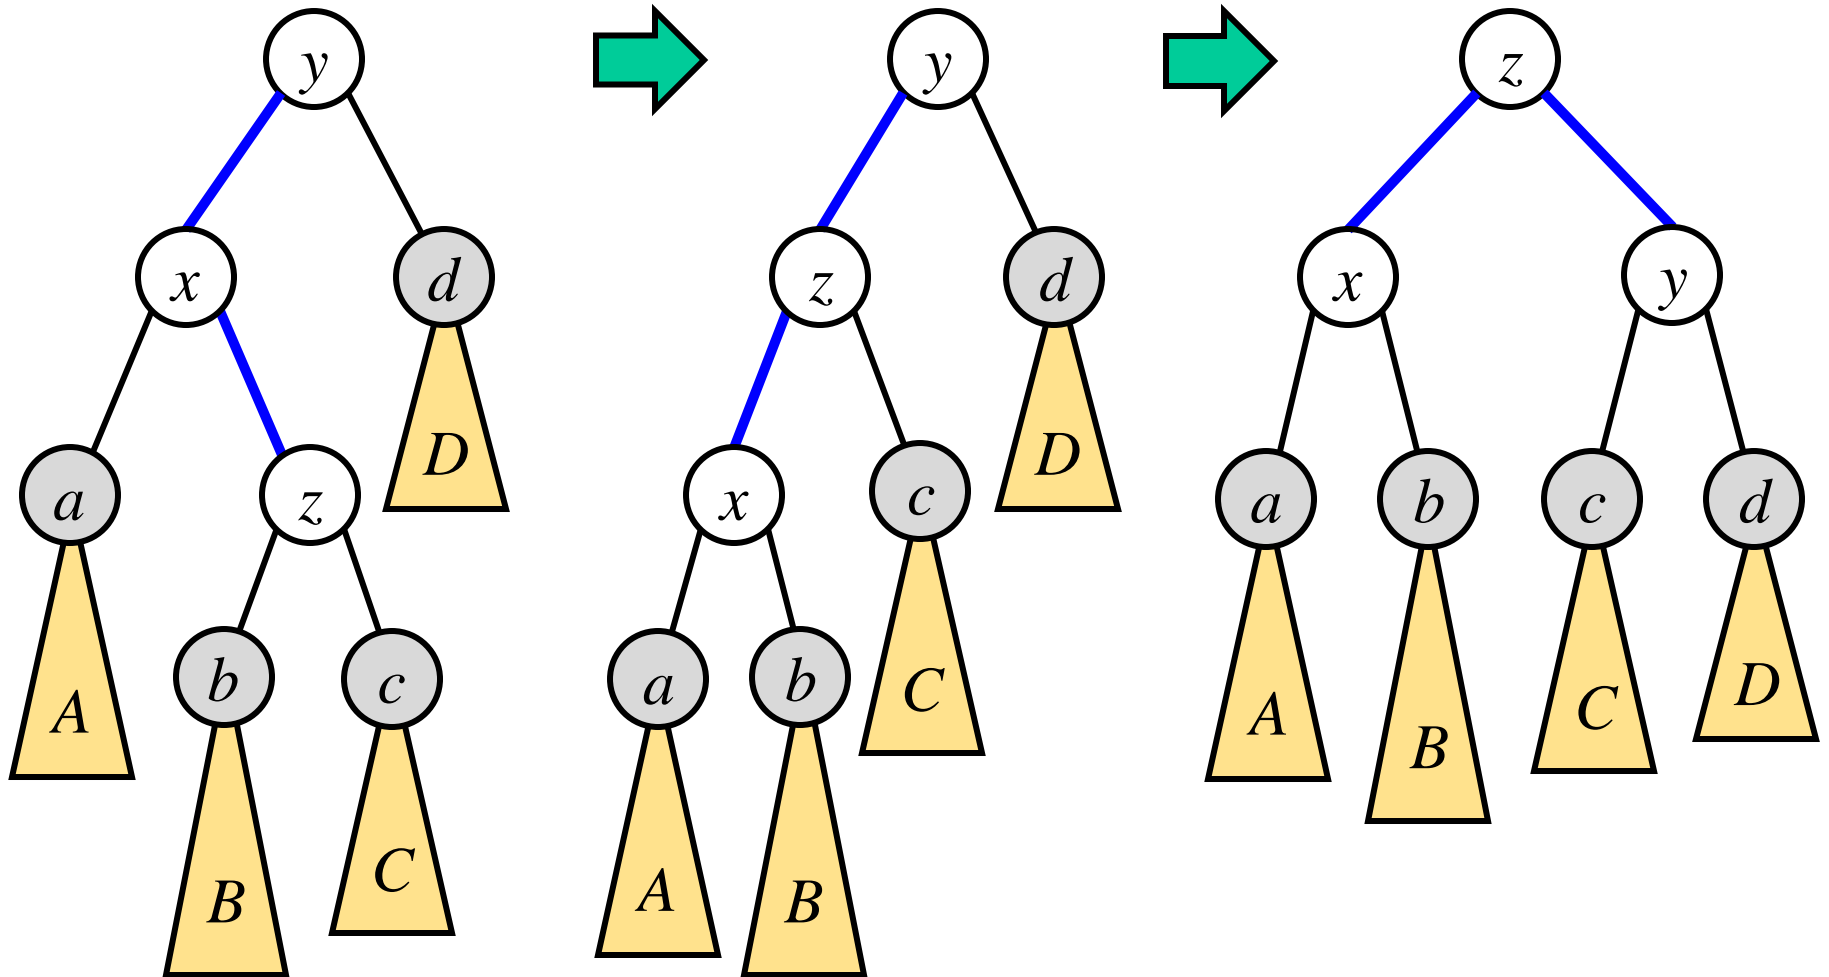
To restructure the tree we use **rotations**

We want to update the tree in $O(\log n)$, or less

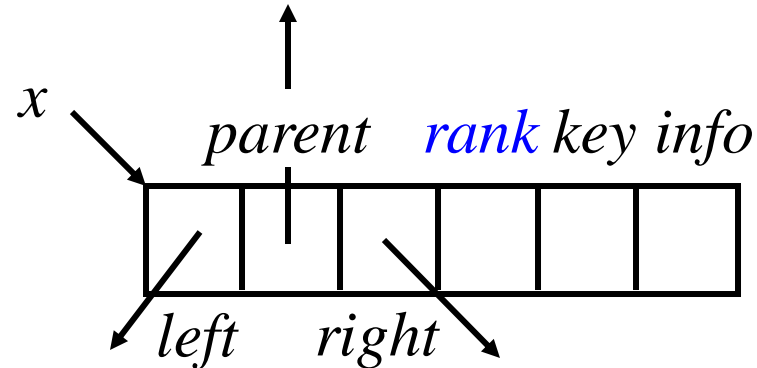
Rotations



Double Rotation



Nodes in AVL trees



Each node has a *rank*

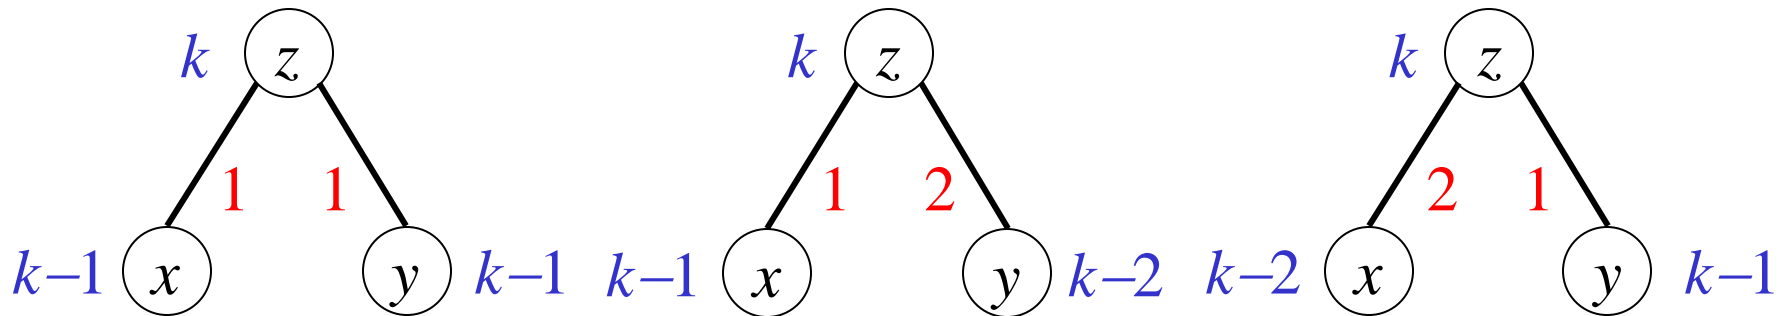
Each edge has a *rank difference*:

$$x.parent.rank - x.rank$$

Before/after each insert/delete *rank = height*

(Enough to keep *rank parity*)

Nodes in AVL trees



Each node has a *rank*

Each edge has a *rank difference*

Before/after each insert/delete *rank = height*

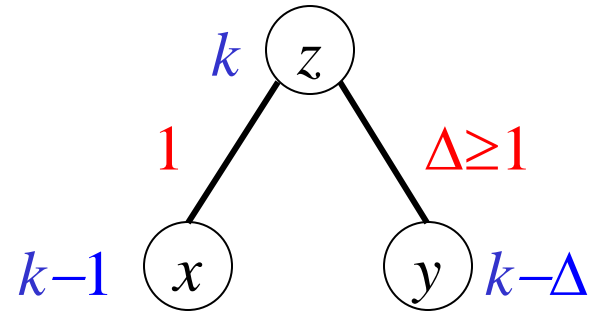
Every node is a 1,1-, 1,2- or 2,1-node

$$\textit{rank} = \textit{height}$$

Lemma: If all leaves have *rank* 0, all *rank differences* are positive, and each parent has a child or *rank difference* 1, then the *rank* of each node is equal to its *height*

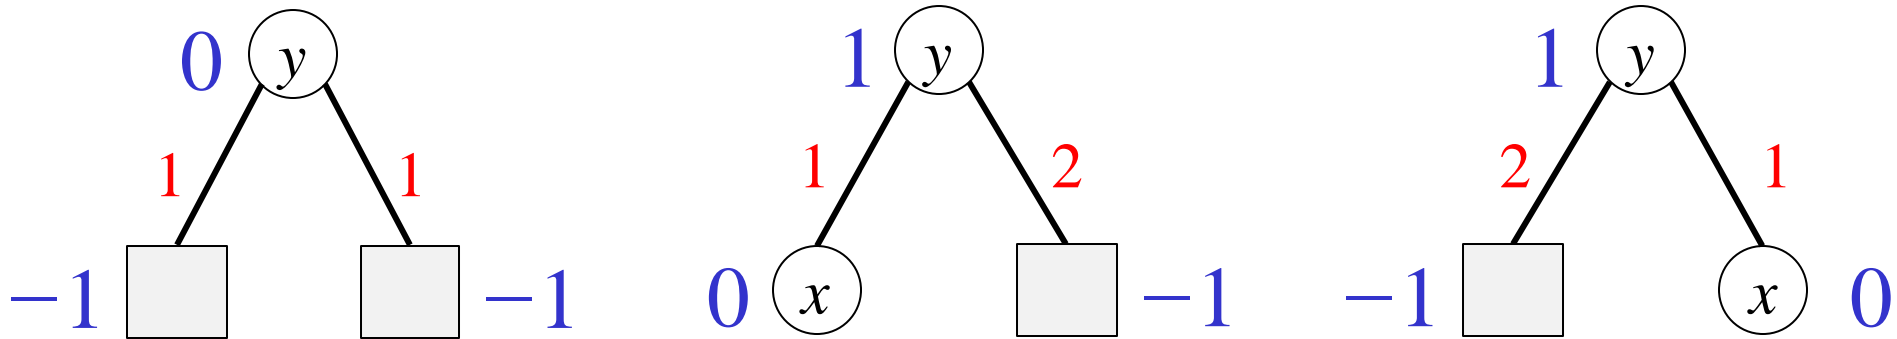
Easy proof by induction:

$$k = 1 + \max\{ k-1, k-\Delta \}$$



Nodes in AVL trees

Internal and external **leaves**



rank of a leaf = 0

rank of a external leaf = -1

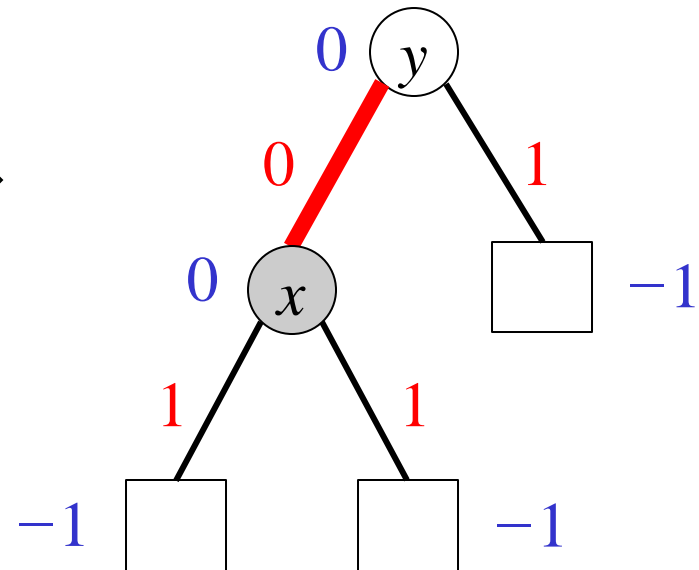
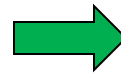
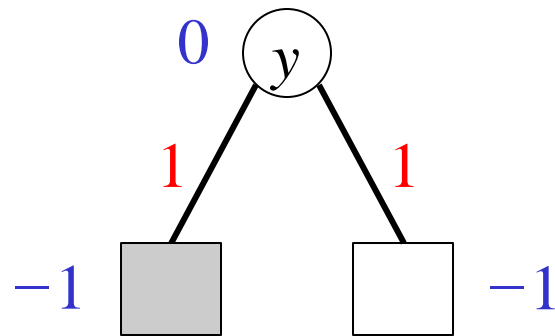
AVL

insertion

AVL Insertion

Replace an external leaf by a new node x

Case A: The parent y is a leaf



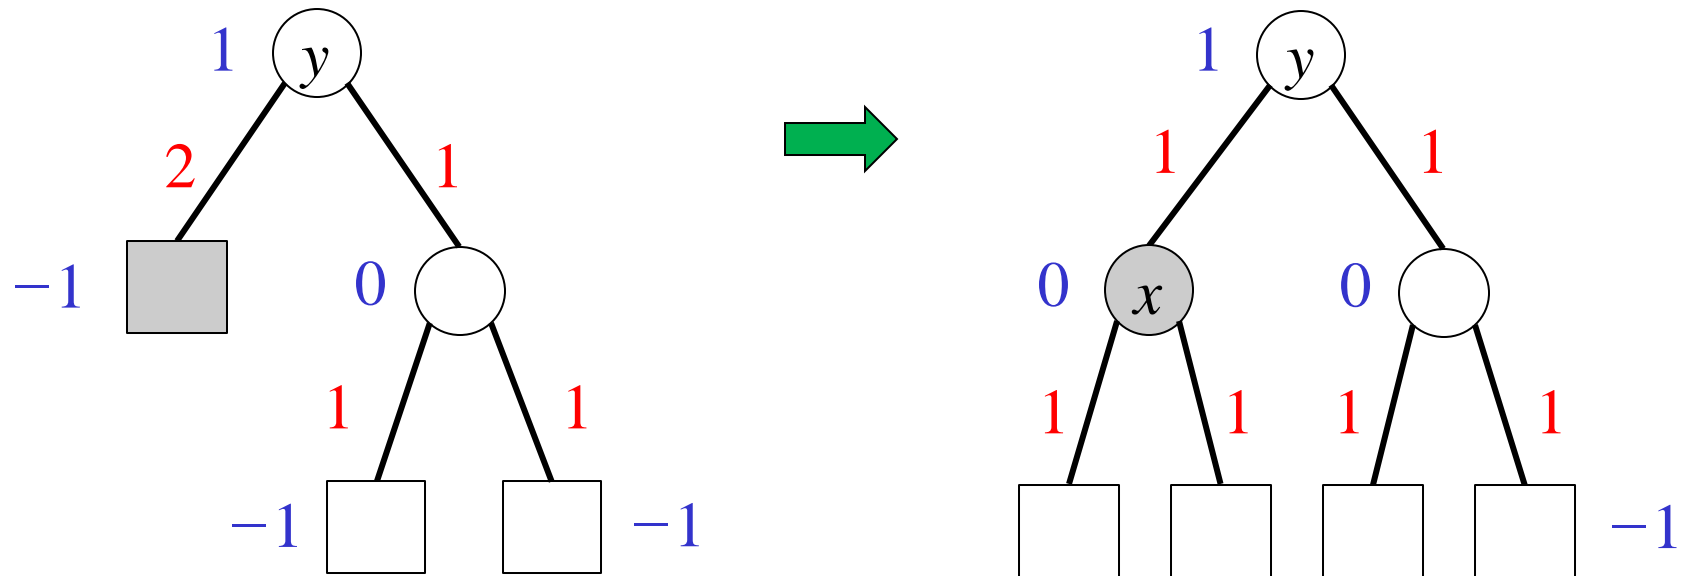
Give x a *rank* of 0

Problem: *rank difference* 0

AVL Insertion

Replace an external leaf by a new node x

Case B: The parent y is not a leaf

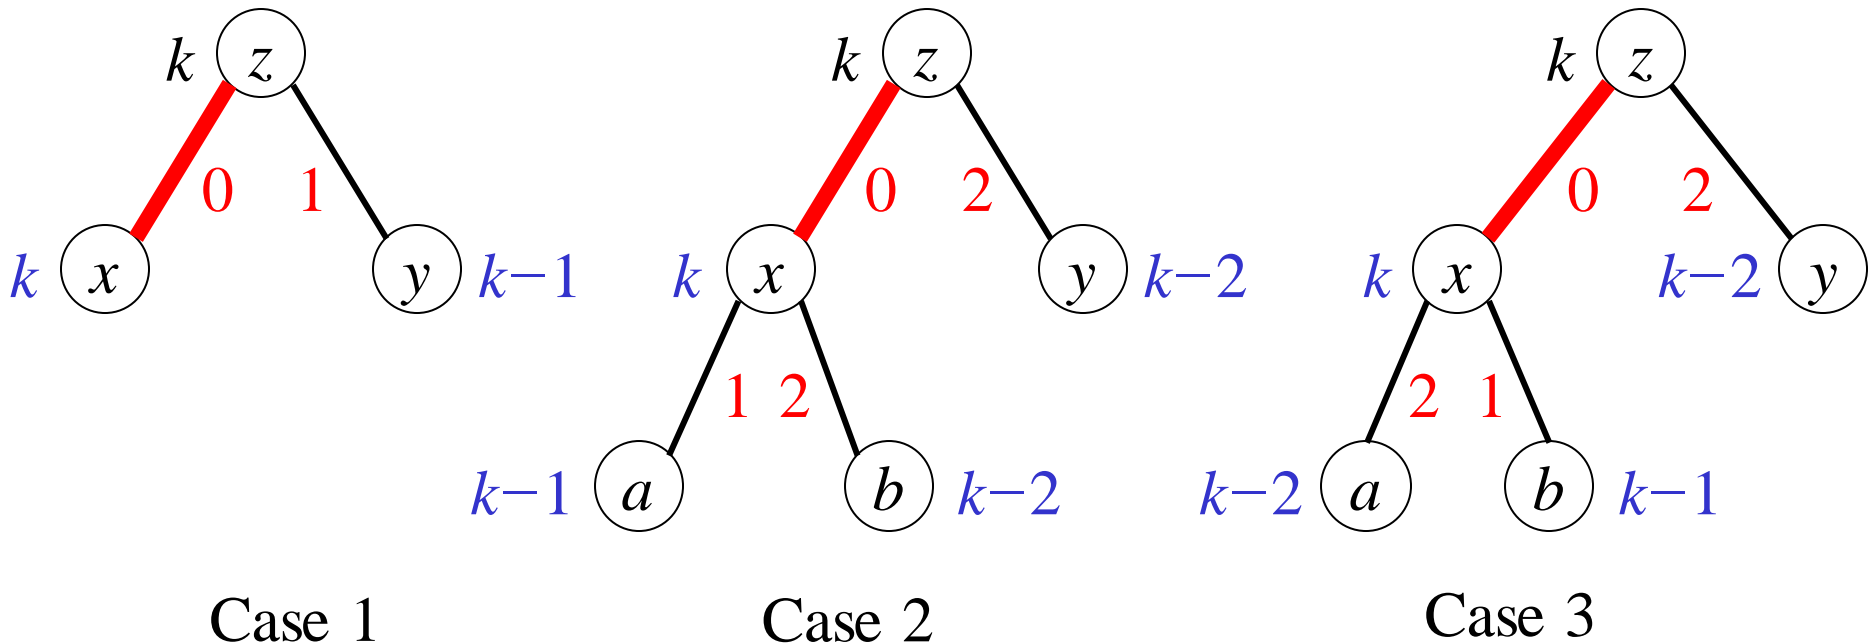


Resulting tree is a valid AVL tree



AVL: Insertion rebalancing

3 cases (up to symmetry)

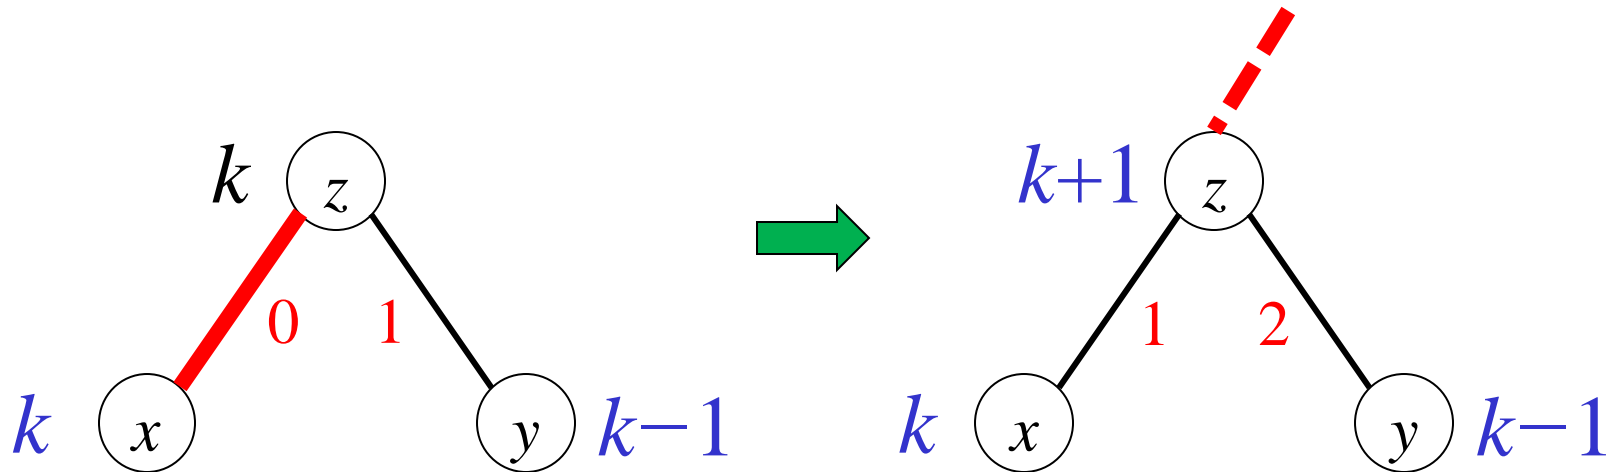


x is the *only* node with *rank difference 0*

In cases 2 and 3, x is a $\{1,2\}$ -node

Rebalancing after insertion

Case 1: Promote

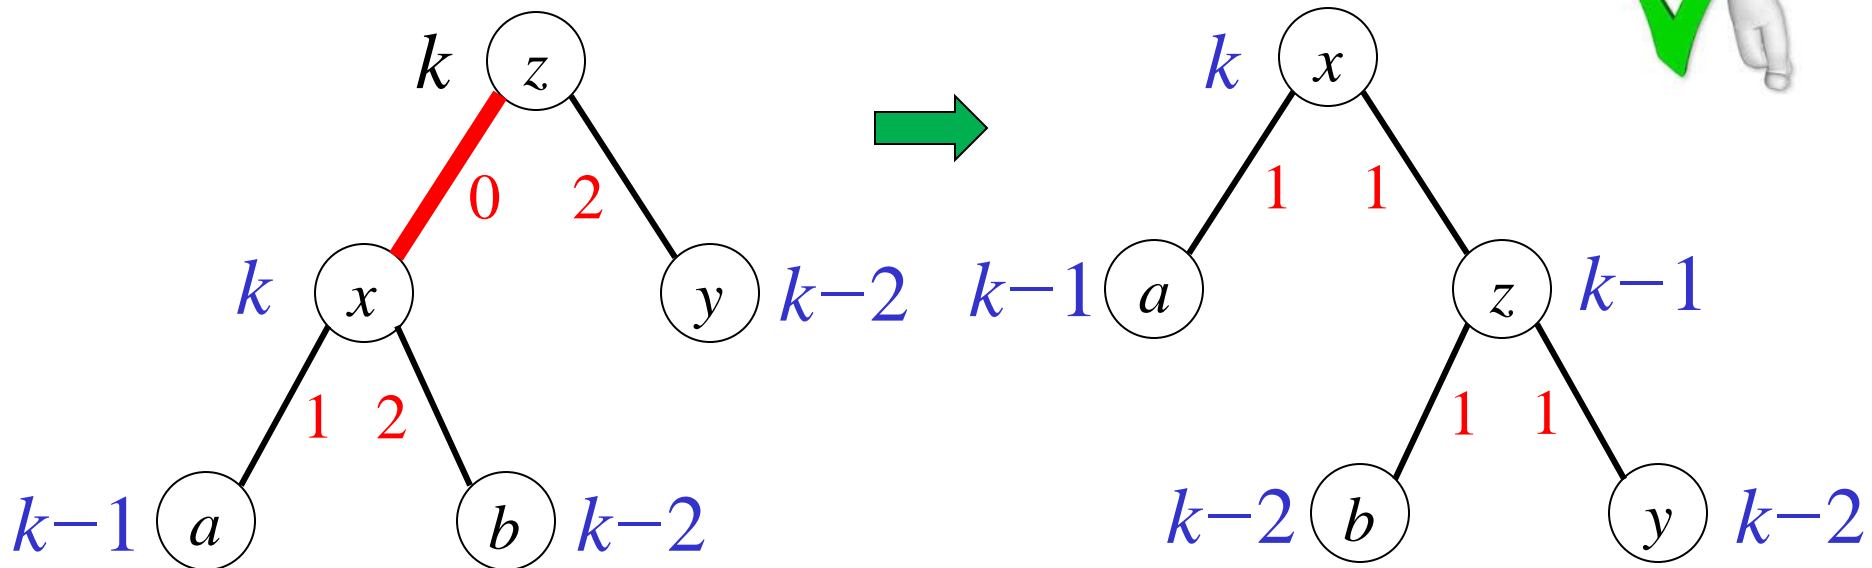


Promote z , i.e., increase its *rank* by 1

Problem is either fixed or moved up

Rebalancing after insertion

Case 2: Single rotation

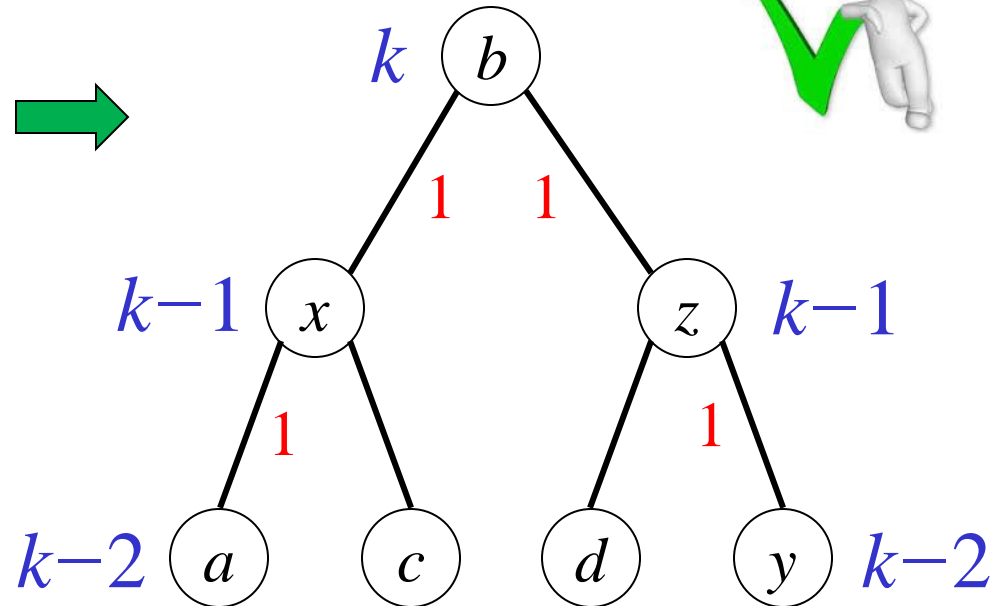
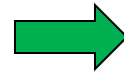
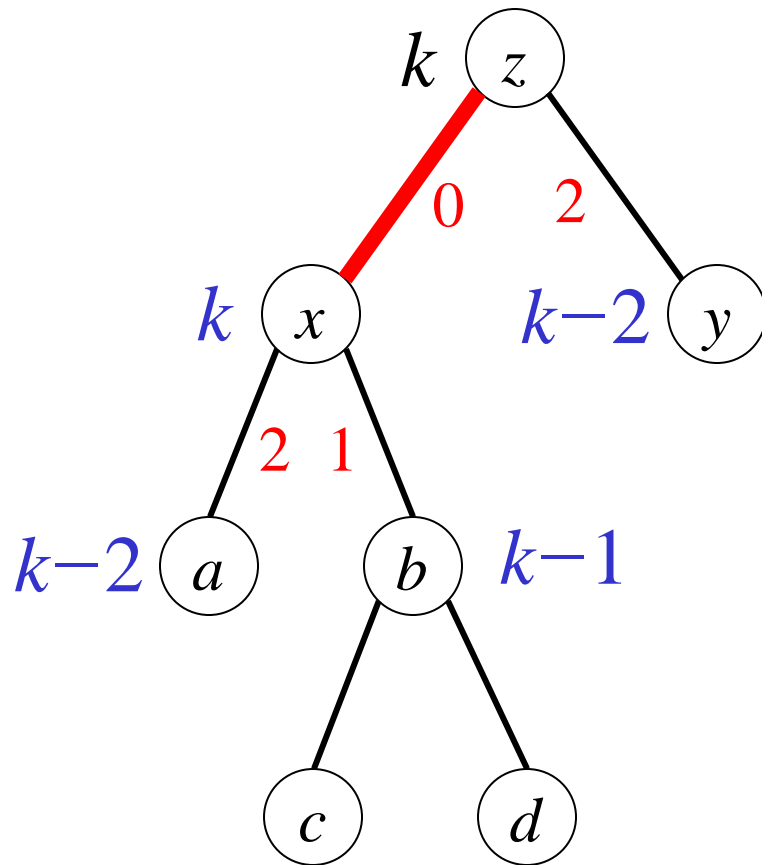


Rotate right Demote z

Rebalancing complete!

Rebalancing after insertion

Case 3: Double rotation



Double rotation

Demote x, z Promote b

Rebalancing complete!

AVL Insertion - Summary

Find insertion point

Insert new node

Rebalance

Number of *promotions* \leq *height* $= O(\log n)$

Number of *rotations* ≤ 2

Worst-case time $= O(\textit{height}) = O(\log n)$

What is the *amortized* number of *rebalancing steps*?

AVL Insertion

Amortized number of *rebalancing steps*

Potential = Φ =
number of 0,1- 1,1-nodes

The insertion itself, and each rebalancing step
change the potential by at most a constant

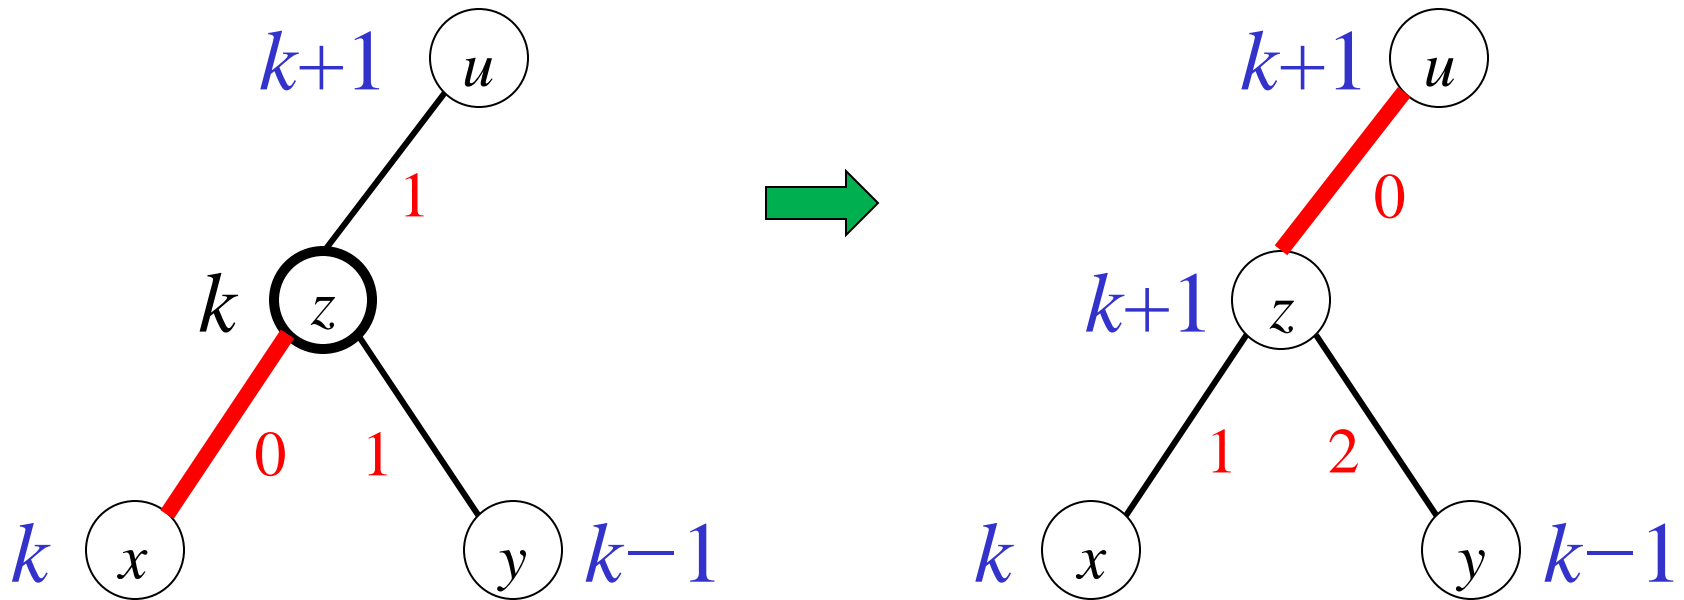
Promotions are the only *non-terminal* cases

Non-terminal promotions *decrease* the potential

$$\textit{amort}(\textit{\#steps}) = O(1)$$

Rebalancing after insertion

Non-terminal promote



Potential of u (and of x, y) does not change

z loses its potential: $\Delta\Phi = -1$

Decrease in potential pays for this step!

AVL deletion

AVL Deletion

Replace item to be deleted with its successor or predecessor, if needed

Delete the appropriate node

Perform a sequence of *demotions*, *rotations* and *double rotations* to restore balance

Somewhat more complicated than insertion

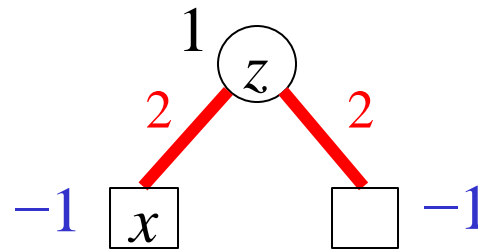
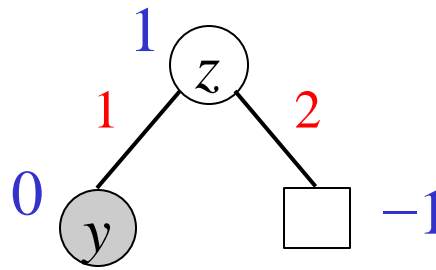
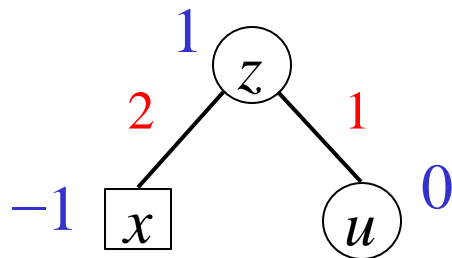
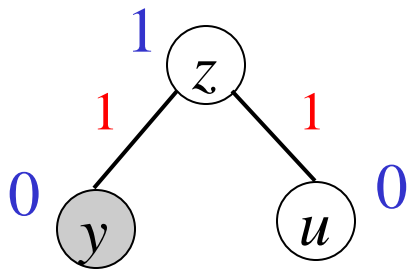
Rotations are not terminal cases

Total deletion time = $O(\text{height}) = O(\log n)$

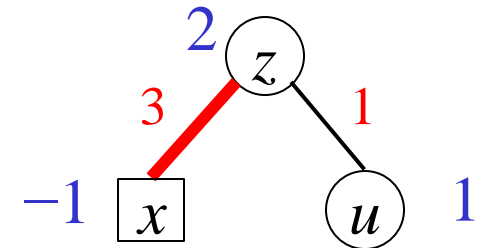
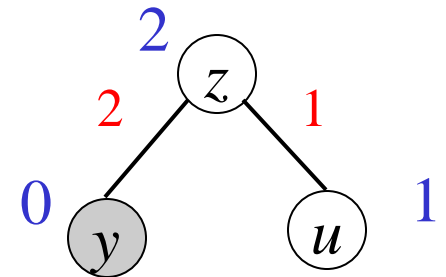
What is the *amortized* cost of rebalancing?

AVL Deletion

Deleting a leaf y



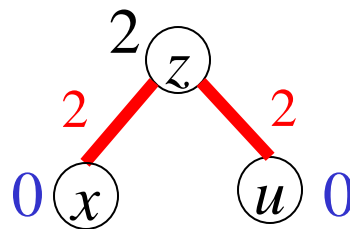
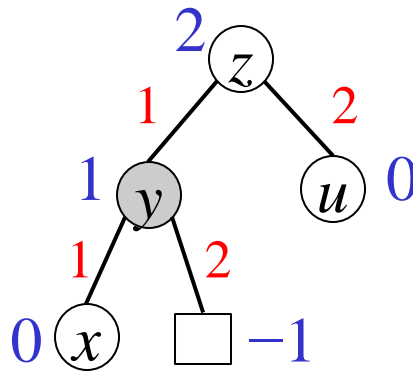
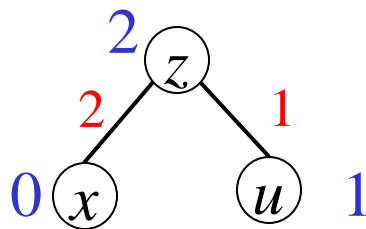
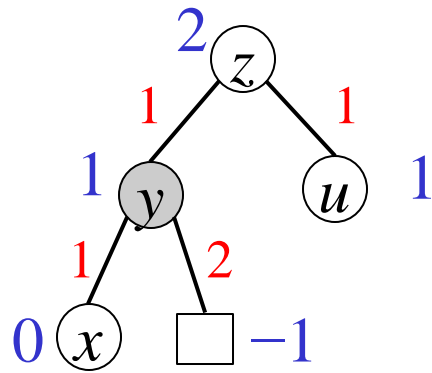
Problem
(Demote z)



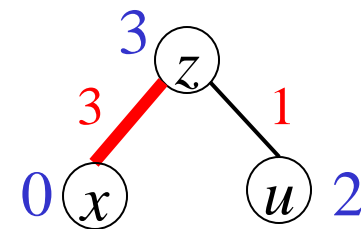
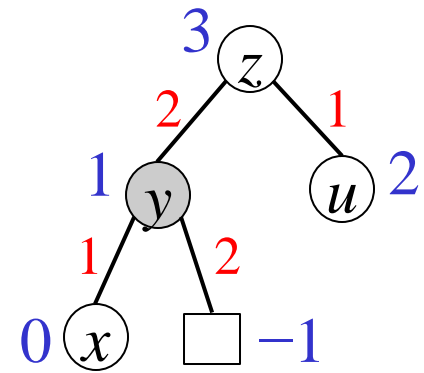
Problem

AVL Deletion

Deleting a unary node y



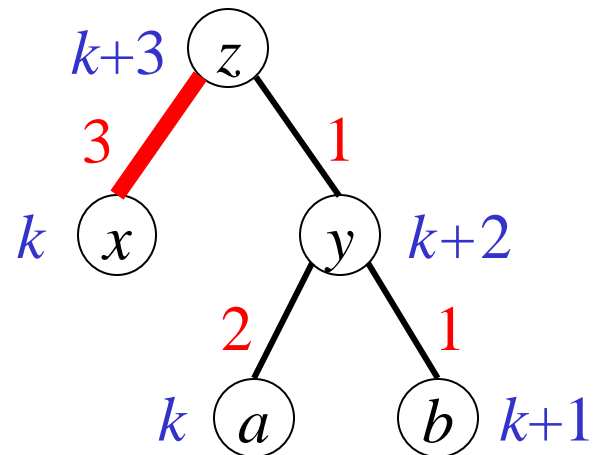
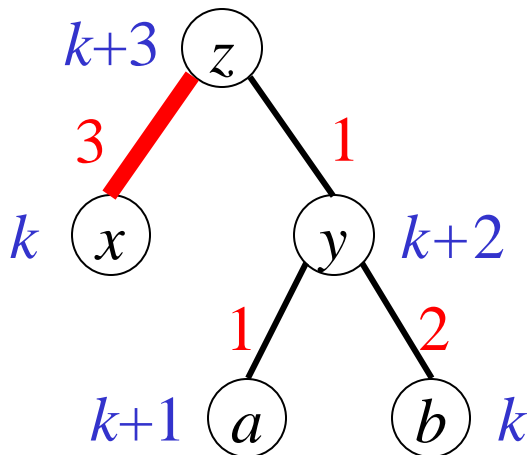
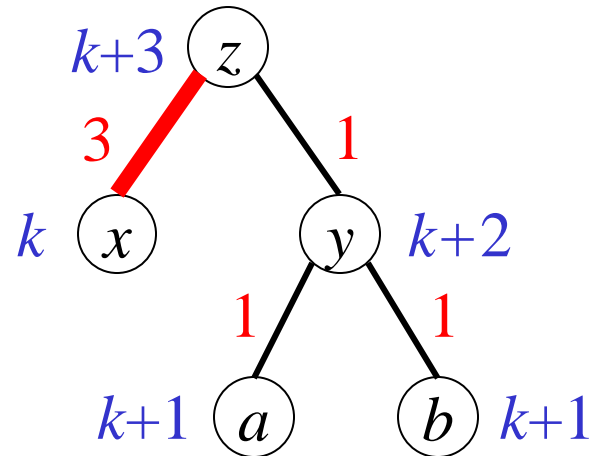
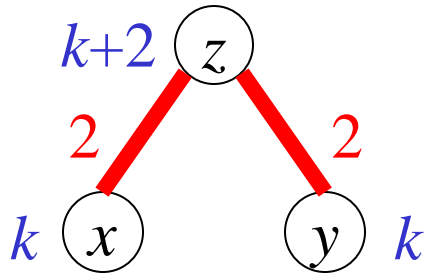
Problem
(Demote z)



Problem

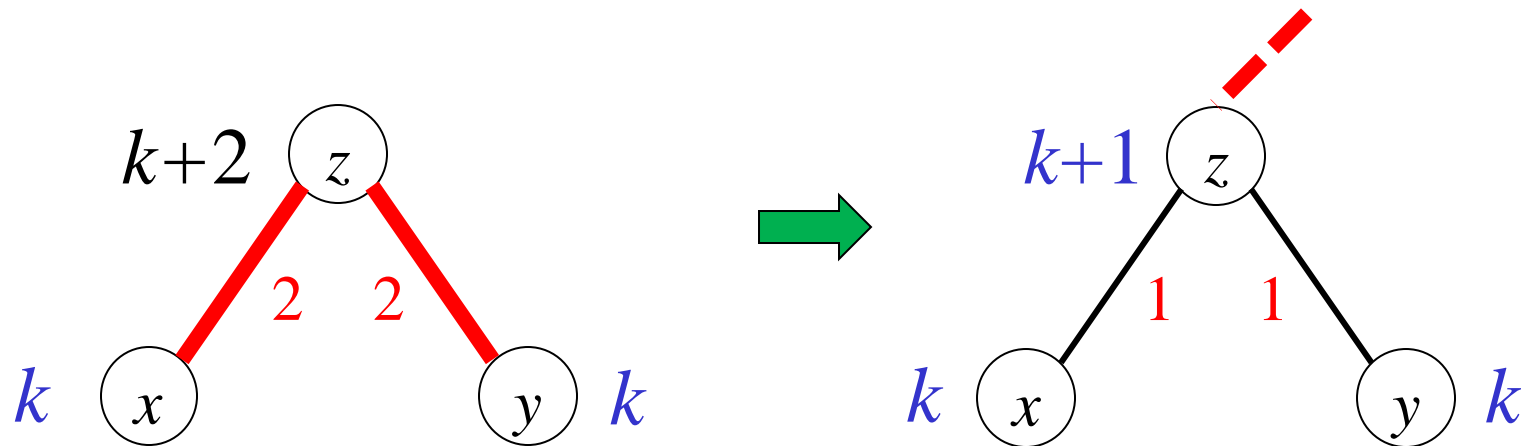
AVL: Deletion rebalancing

4 cases (up to symmetry)



AVL: Rebalancing after deletion

Case 1: Demote

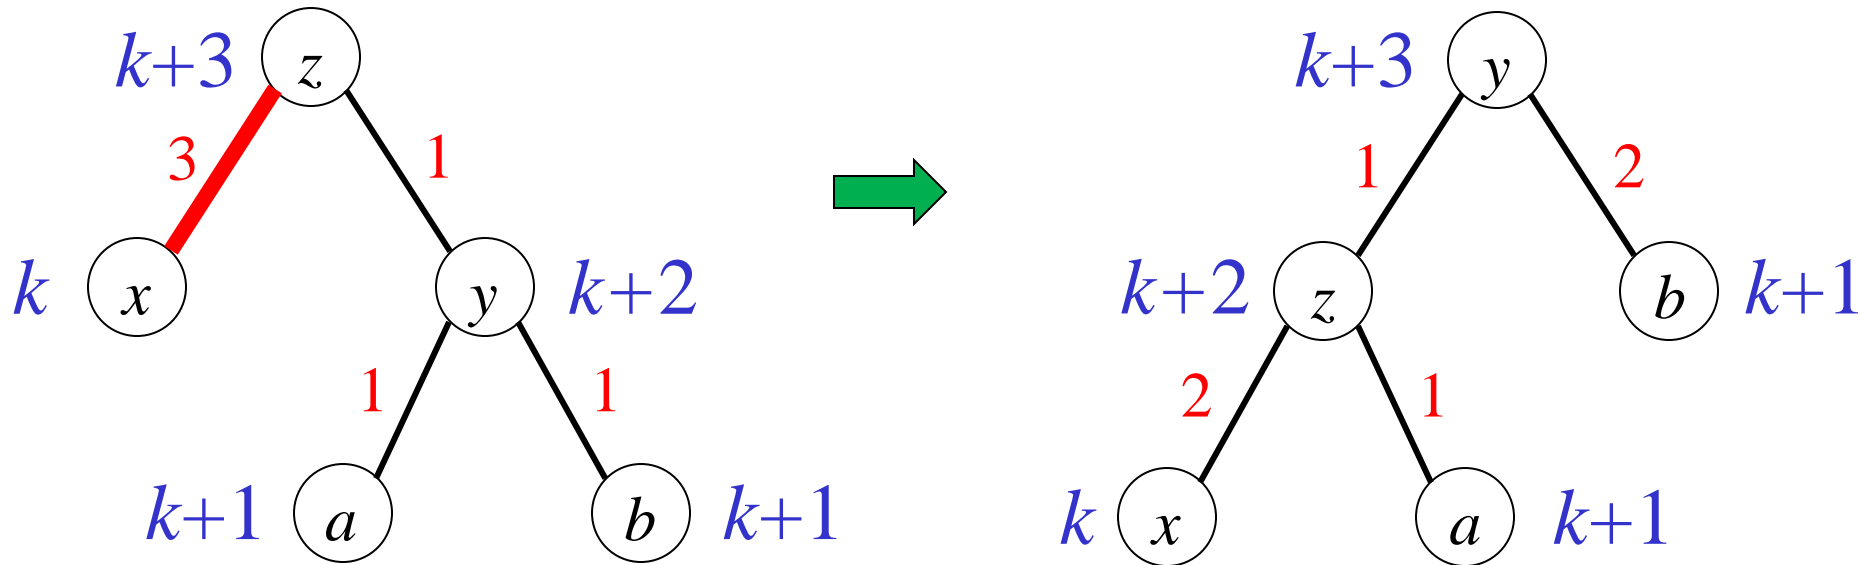


Demote z , i.e., decrease its *rank* by 1

Problem is either fixed or moved up

AVL: Rebalancing after deletion

Case 2: Single Rotation (a)



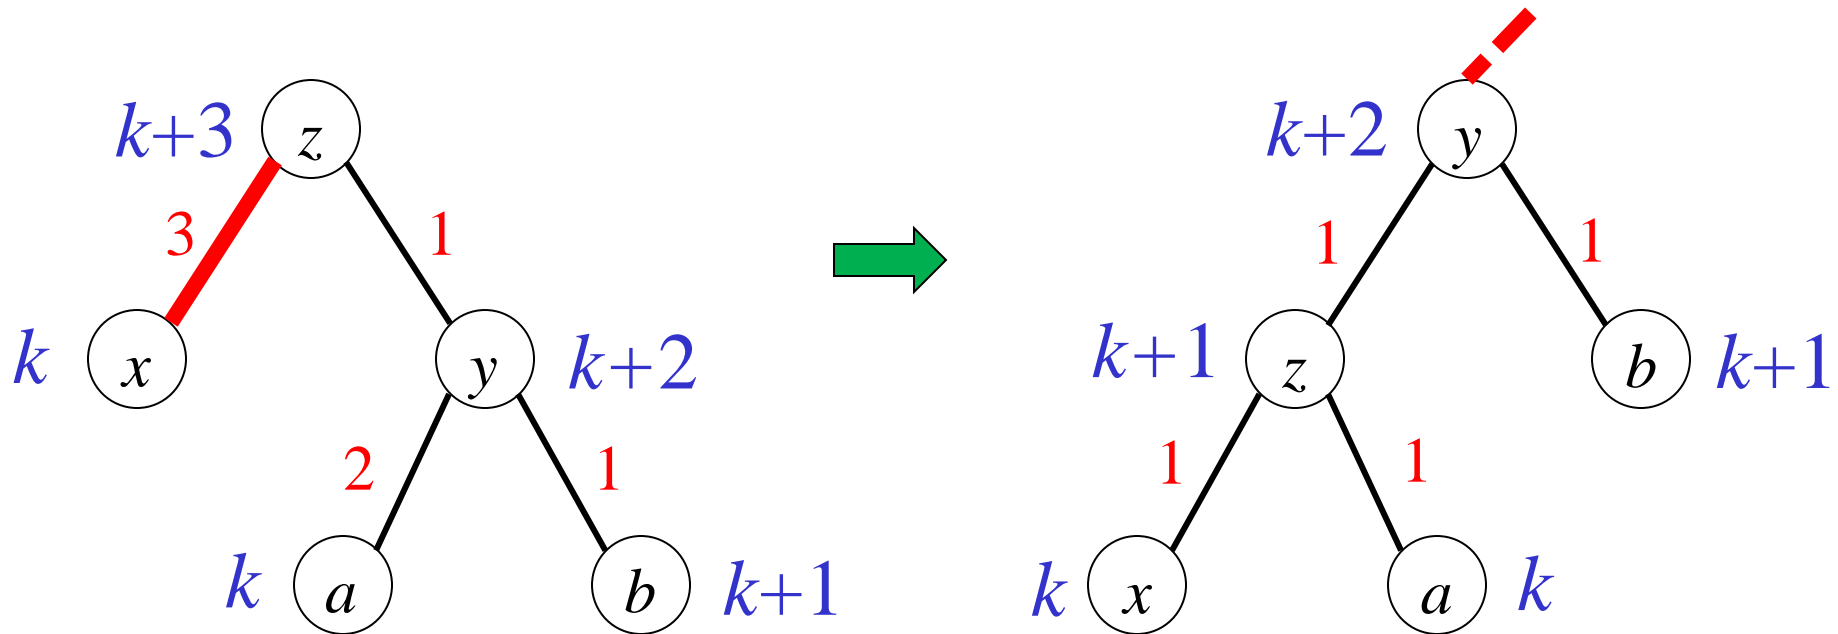
Rotate left

Demote z Promote y



AVL: Rebalancing after deletion

Case 3: Single Rotation (b)

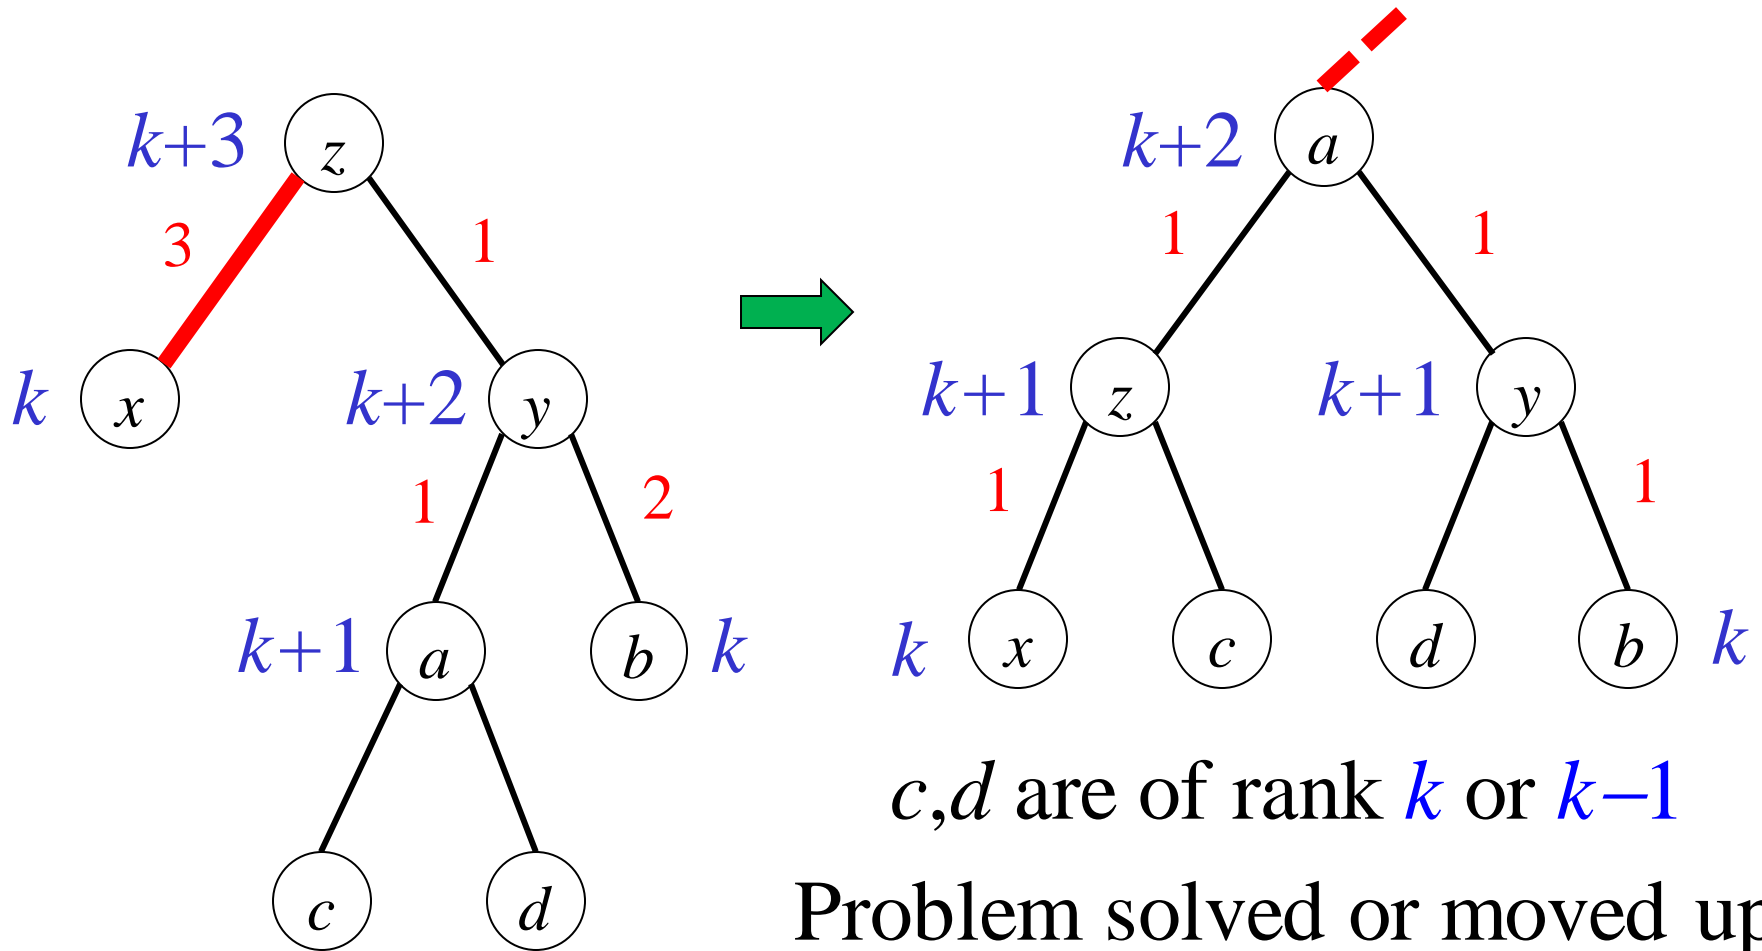


Rotate left
Demote z twice

Problem solved or moved up

AVL: Rebalancing after deletion

Case 4: Double Rotation



AVL Deletion - Summary

Replace item to be deleted with its successor or predecessor, if needed

Delete the appropriate node

Perform a sequence of *demotions*, *rotations* and *double rotations*

Somewhat more complicated than insertion

Rotations are not terminal cases

Total deletion time = $O(\text{height}) = O(\log n)$

What is the *amortized* cost of *rebalancing*?

AVL Trees – Cost of *rebalancing*

Worst-case cost of *rebalancing*, after both **insertions** and **deletions**, is $O(\log n)$

If there are only **insertions**, the *amortized* cost of *rebalancing*, as we saw, is $O(1)$

If there are only **insertions**, and then only **deletions**, the *amortized* cost of *rebalancing* is again $O(1)$

But, if **insertions** and **deletions** are intermixed, the *amortized* cost of *rebalancing* may be $\Omega(\log n)$

Can the *amortized* cost of *rebalancing*, in the general case, be brought down to $O(1)$?

WAVL trees

[Haeupler-Sen-Tarjan (2009)]

At most two *rotations*
both in insertions and deletions

Amortized cost of *rebalancing* is $O(1)$

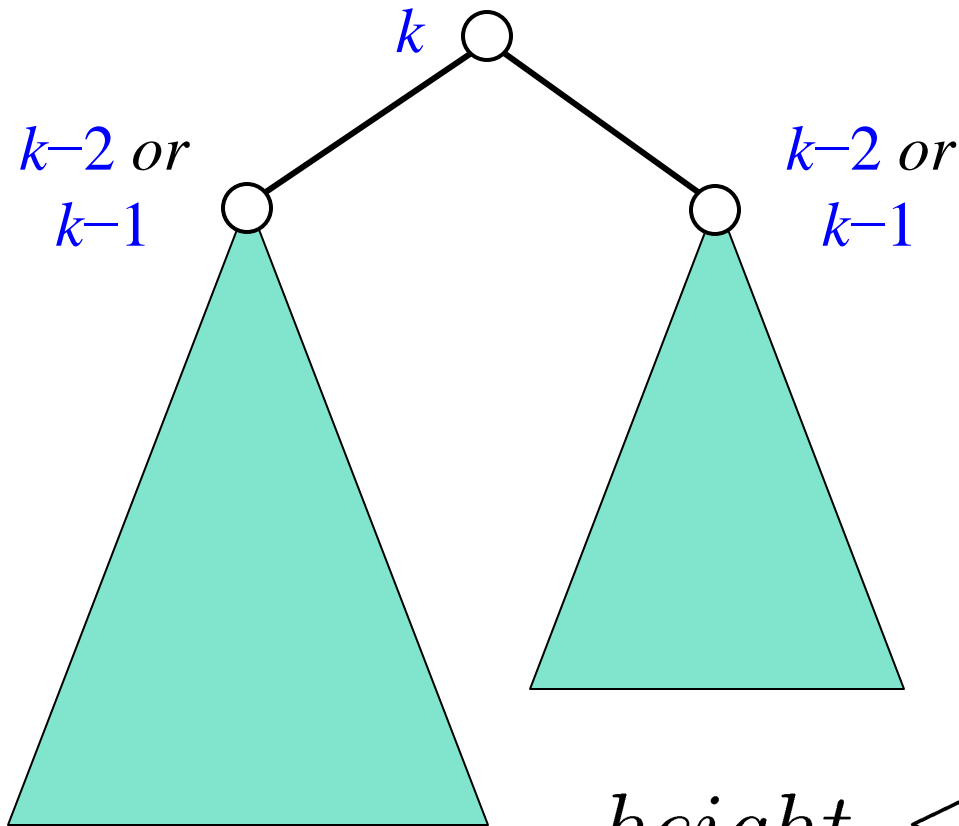
Allow *2,2*-nodes

Every (internal) leaf is still a *1,1*-node

rank \neq *height*

WAVL trees

[Haeupler-Sen-Tarjan (2009)]



WAVL trees are
search trees

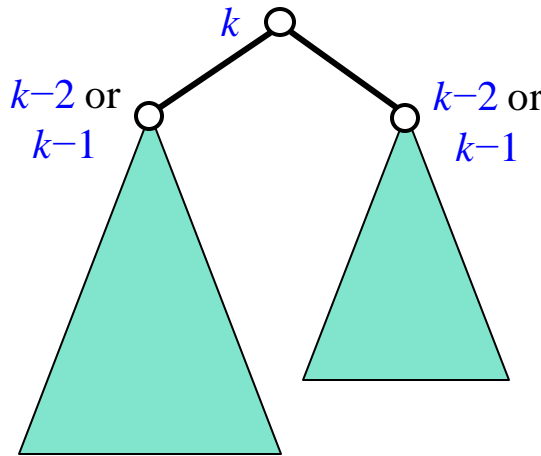
Each *rank-difference*
is 1 or 2

rank of each
(internal) leaf is 0

$$height \leq rank \leq 2 height$$

WAVL trees

[Haeupler-Sen-Tarjan (2009)]



S_k – minimal number of nodes in an WAVL tree of *rank* k

$$S_k = 2S_{k-2} + 1$$

$$S_{-1} = 0 \quad , \quad S_0 = 1$$

By induction: $S_k \geq 2^{\lceil k/2 \rceil}$

$$height \leq rank \leq 2 \log_2 n$$

WAVL Insertion

Exactly like AVL insertion

No 2,2-nodes created

2,2-nodes may be destroyed

(Check this!)

If there are only insertions,
WAVL trees are just AVL trees

WAVL deletion

WAVL Deletion

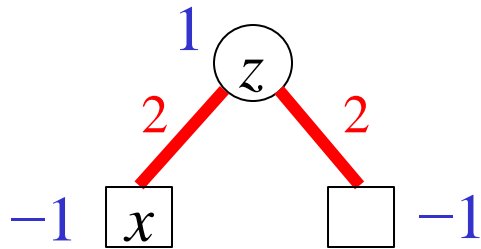
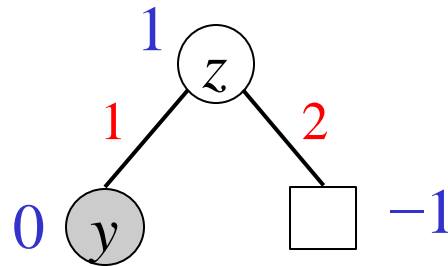
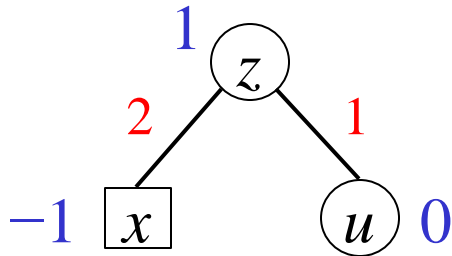
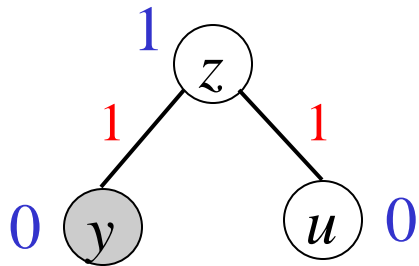
New cases to consider

Deletion becomes somewhat simpler

Rotations are now terminal cases

WAVL Deletion

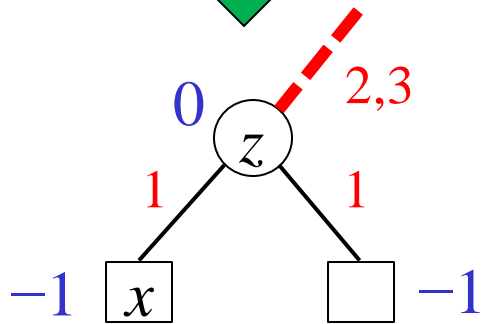
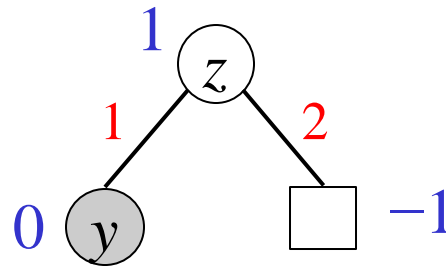
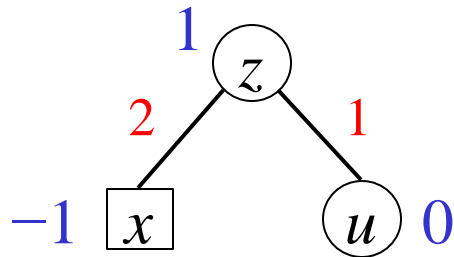
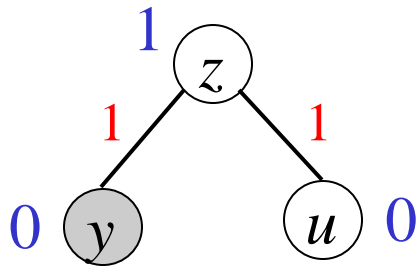
Deleting a leaf y



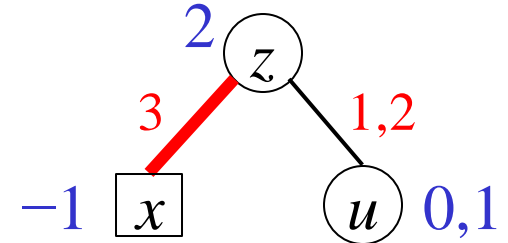
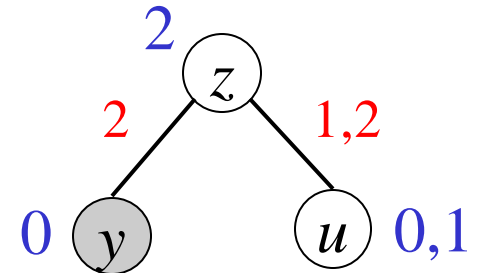
Problem! (Why?)
(Demote z)

WAVL Deletion

Deleting a leaf y



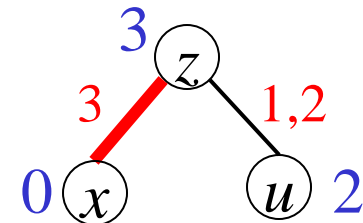
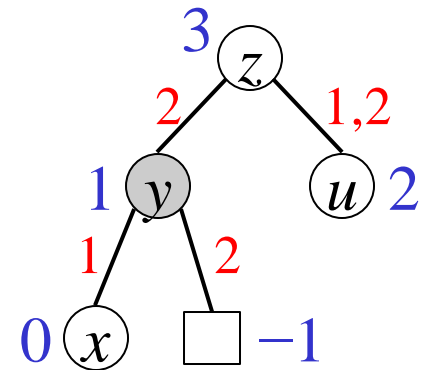
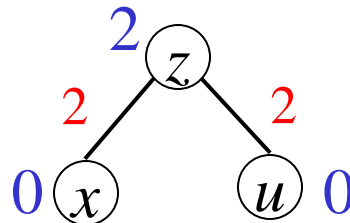
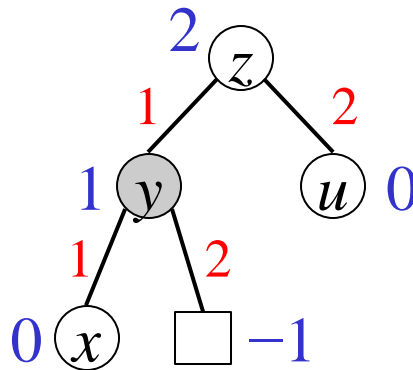
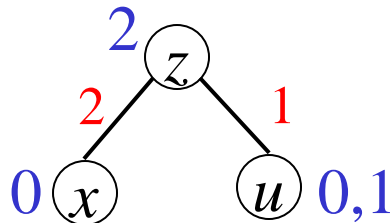
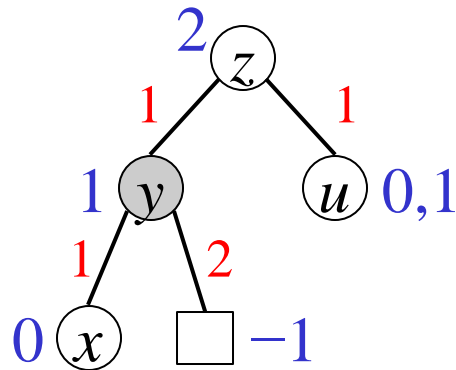
Demote z
(May cause a problem)



Problem

WAVL Deletion

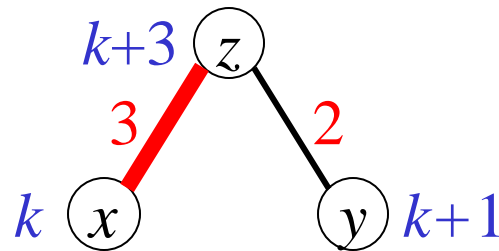
Deleting a unary node y



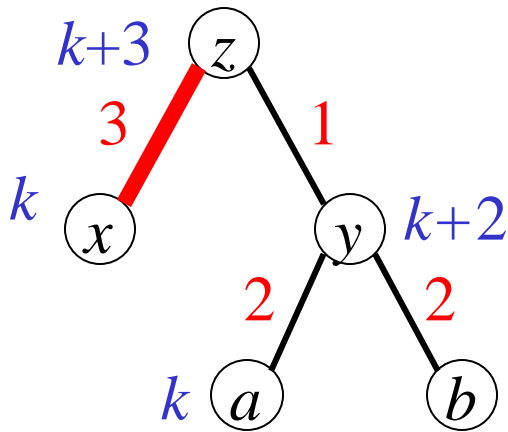
Problem

WAVL: Deletion rebalancing cases

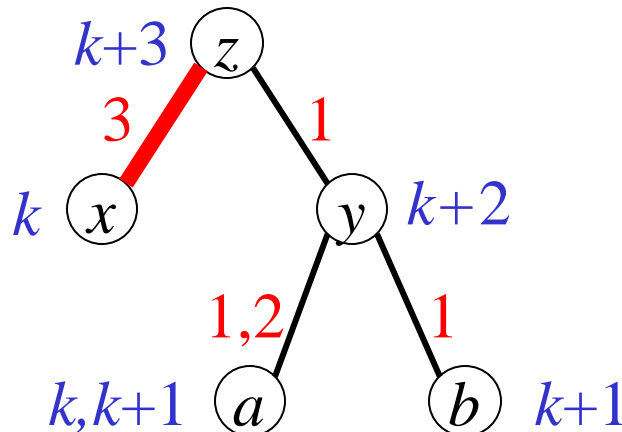
4 cases (up to symmetry)



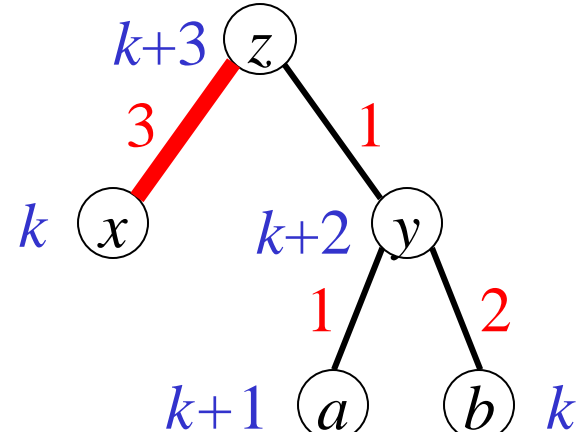
Case 1: Demote



Case 2: Double Demote



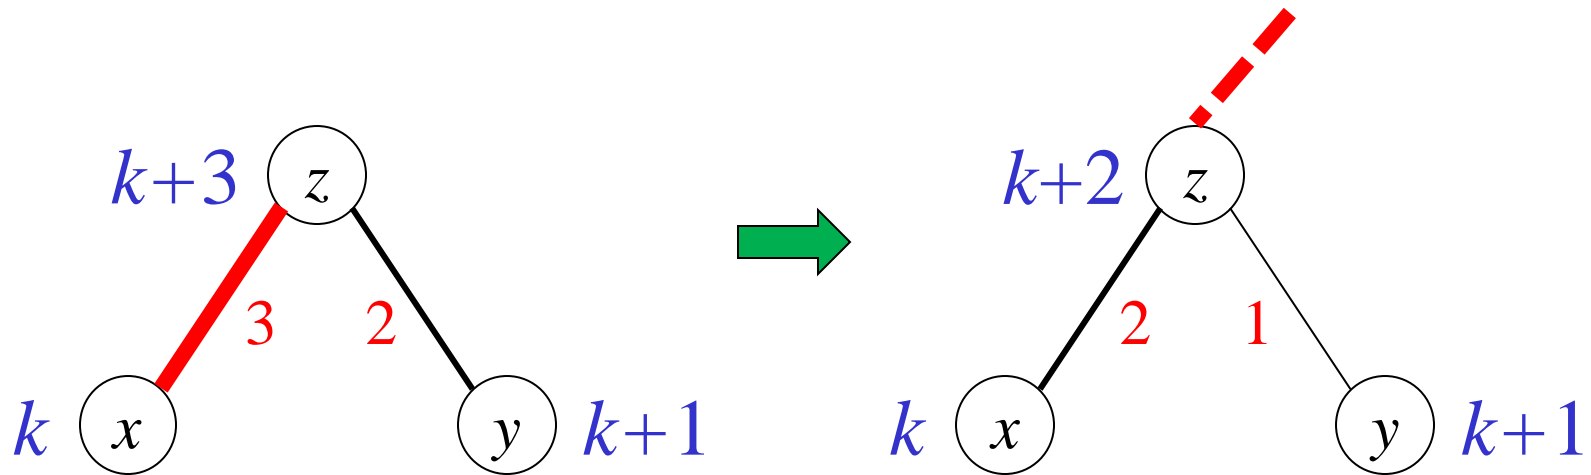
Case 3: Rotate



Case 4: Double Rotate

WAVL: Rebalancing after deletion

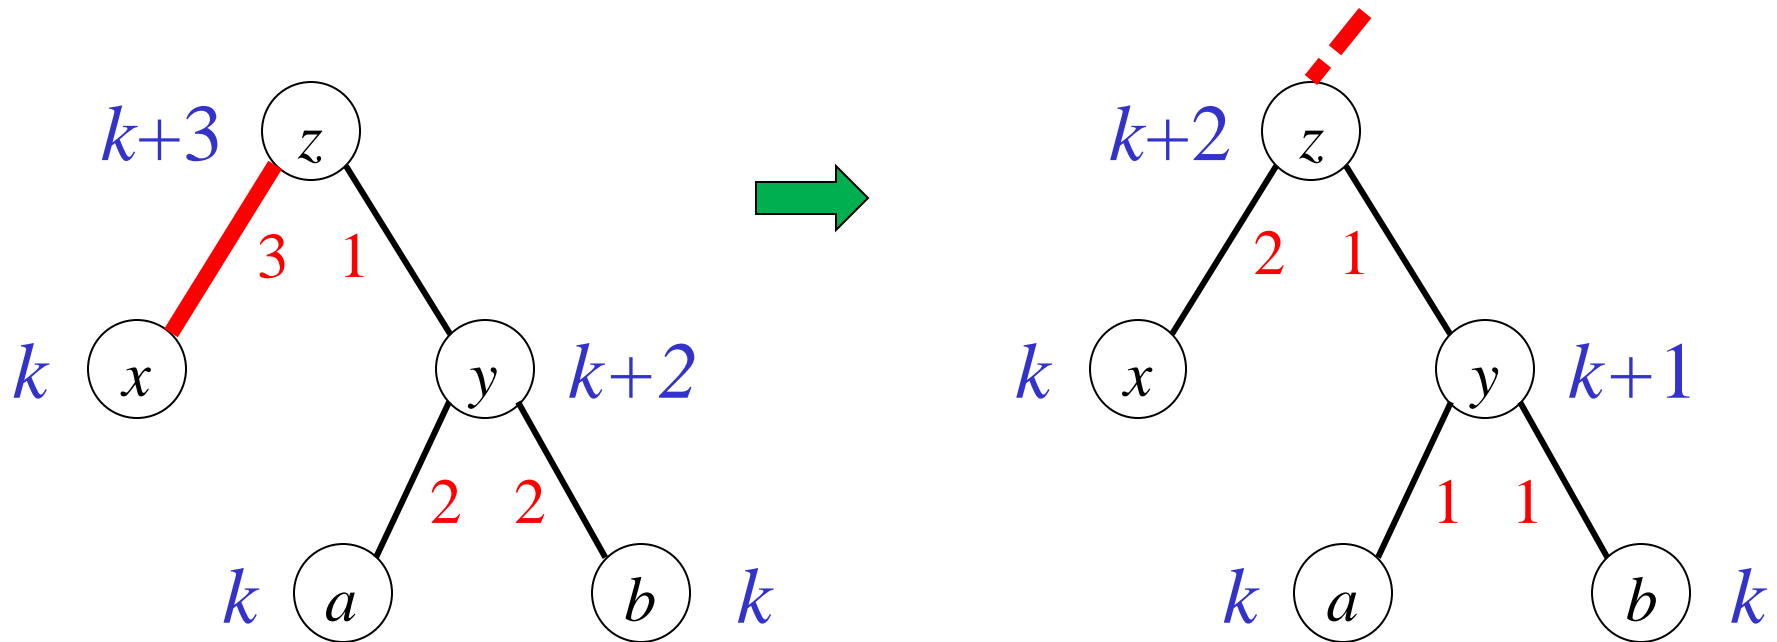
Case 1: Demote



Problem is either fixed or moved up

WAVL: Rebalancing after deletion

Case 2: Double demote

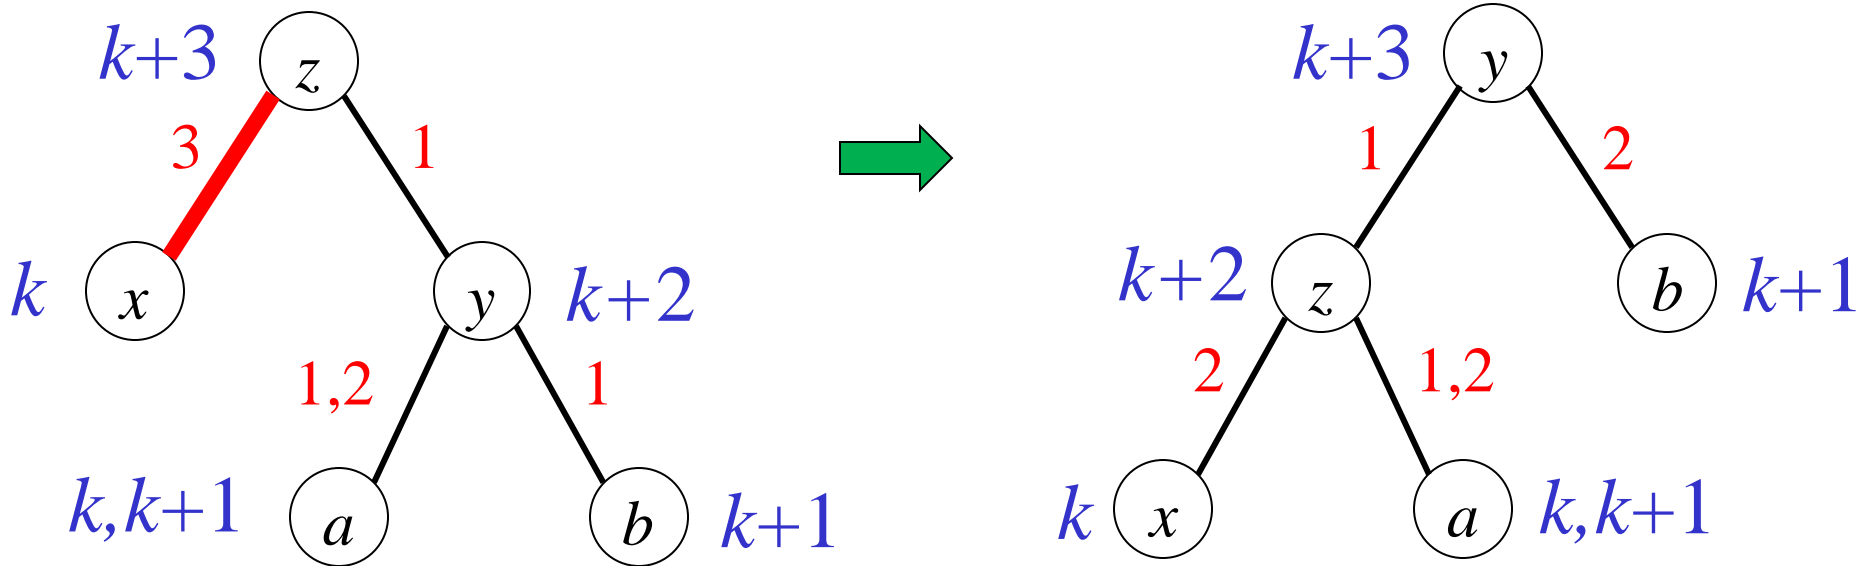


Demote z and y

Problem either solved or moved up

WAVL: Rebalancing after deletion

Case 3: Rotate

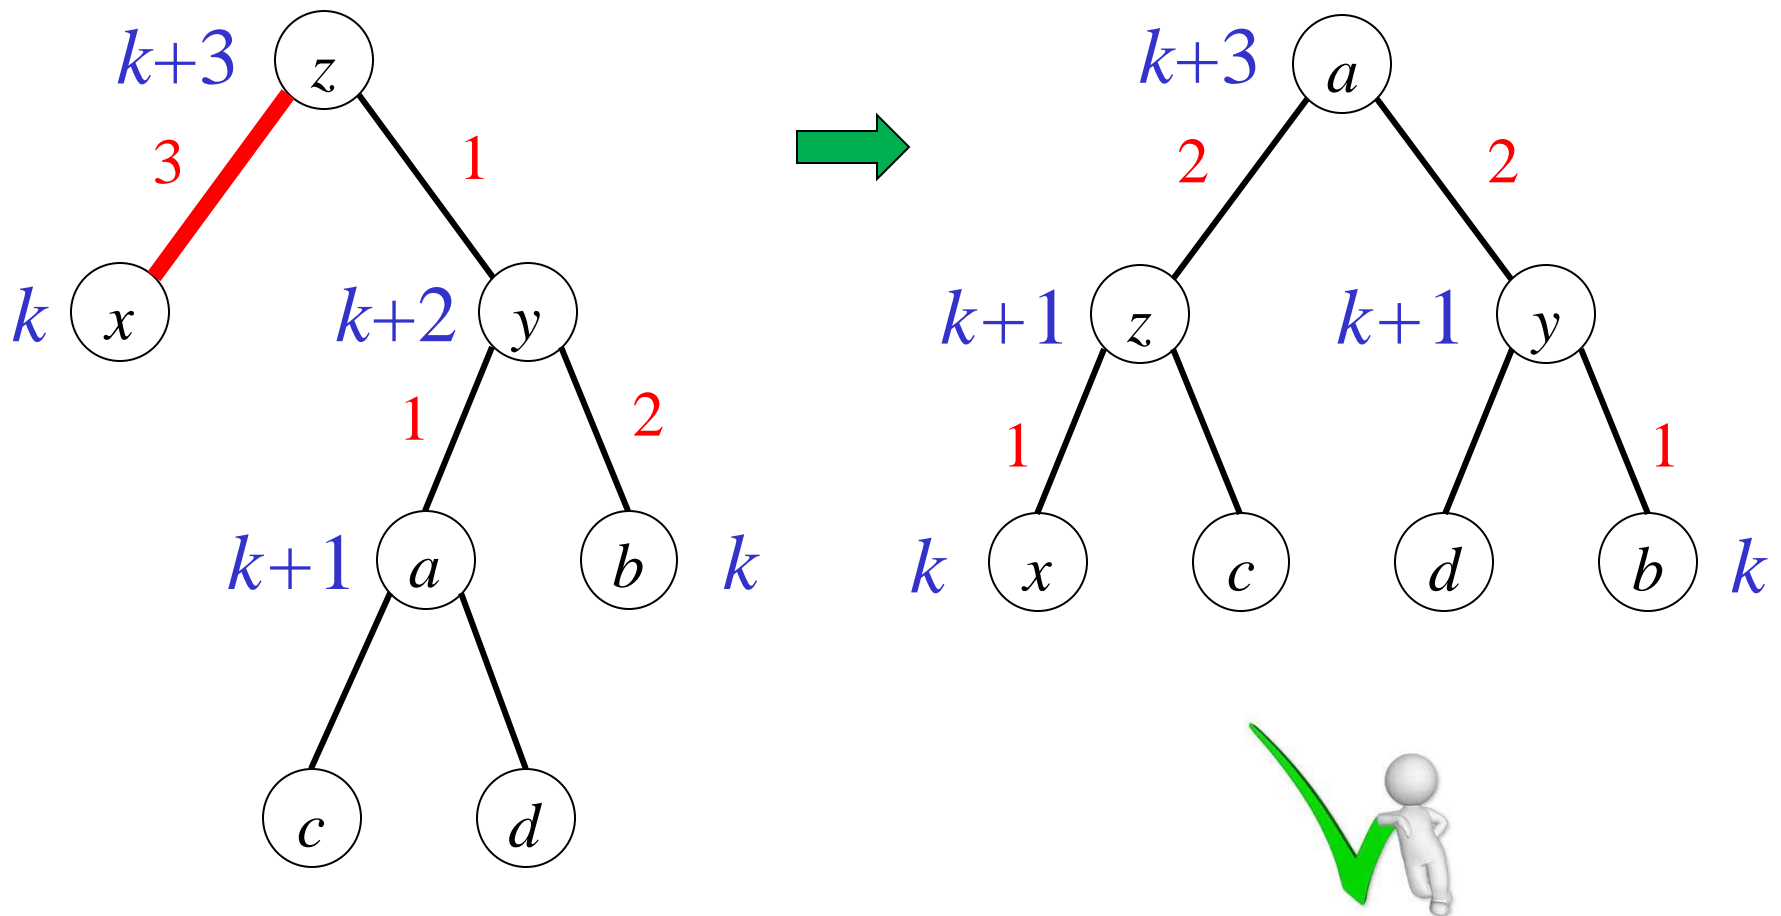


If z is a 2,2-leaf, demote it



WAVL: Rebalancing after deletion

Case 4: Double Rotate



WAVL Insertion/Deletion - Summary

promotions/demotions \leq *height* $= O(\log n)$

Number of *rotations* ≤ 2

Worst-case time $= O(\textit{height}) = O(\log n)$

What is the *amortized* number
of *rebalancing steps*?

WAVL Insertion/Deletion

Amortized number of *balancing steps*

$$\begin{aligned} \text{Potential} = \Phi = \\ & (\text{number of } 0,1\text{- and } 1,1\text{-nodes}) \\ & + 2 \times (\text{number of } 3,2\text{- } 2,2\text{-nodes}) \end{aligned}$$

Insertions/Deletions themselves, and each rebalancing step, change the potential by at most a constant

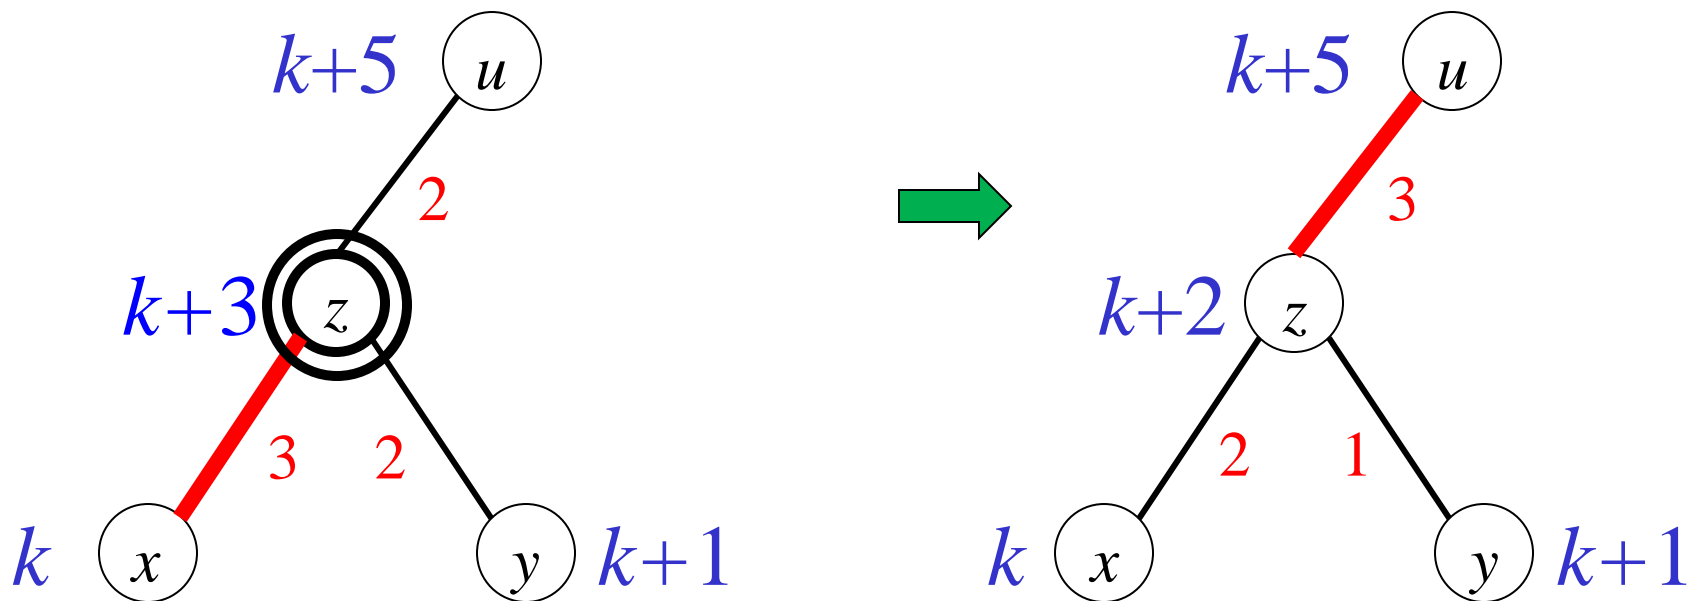
Promotions/Demotions are the only *non-terminal* steps

Non-terminal steps *decrease* the potential

$$\text{amort}(\# \text{steps}) = O(1)$$

WAVL: Rebalancing after deletion

Non-terminal Demote

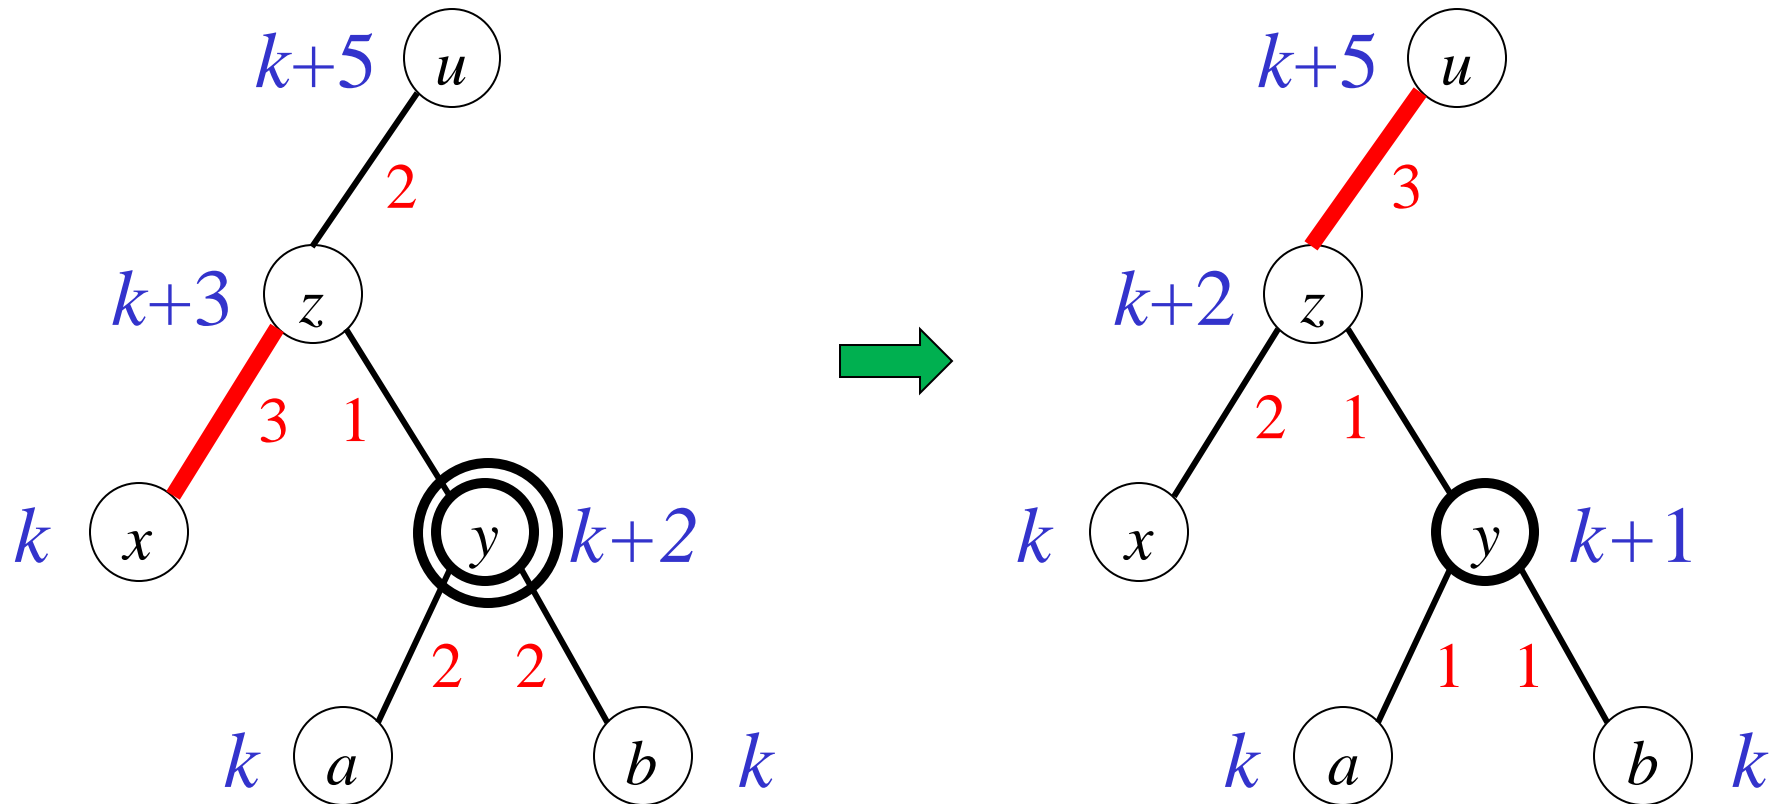


Potential of u (and of x, y) does not change

Potential of z drops from 2 to 0: $\Delta\Phi = -2$

WAVL: Rebalancing after deletion

Non-terminal Double Demote



Potential of y drops from 2 to 1: $\Delta\Phi = -1$

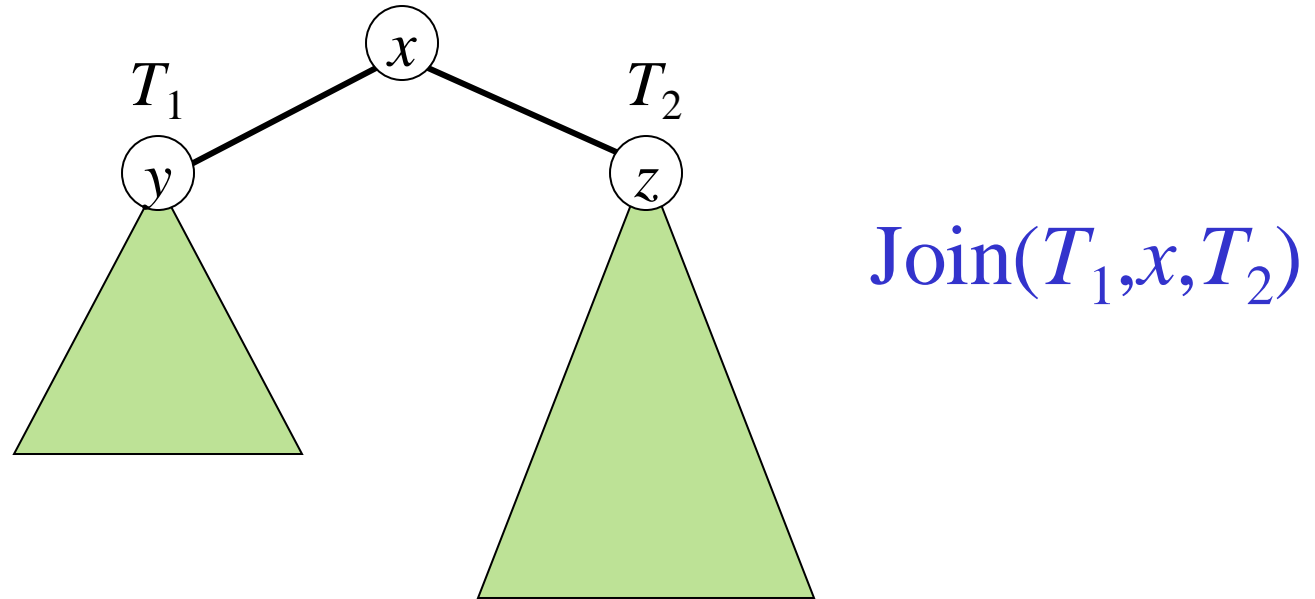
AVL vs. WAVL

	Depth	Number of Rotations per update	Amortized cost of rebalancing
AVL	$1.45 \log_2 n$	$O(\log n)$	$O(\log n)$
WAVL	$2 \log_2 n$	2	$O(1)$

Theorem: The depth of a WAVL tree generated by a sequence of m insertions, and an arbitrary number of deletions, is at most $\log_\phi m \leq 1.45 \log_2 m$

Joining and Splitting binary search trees

Joining two binary search trees



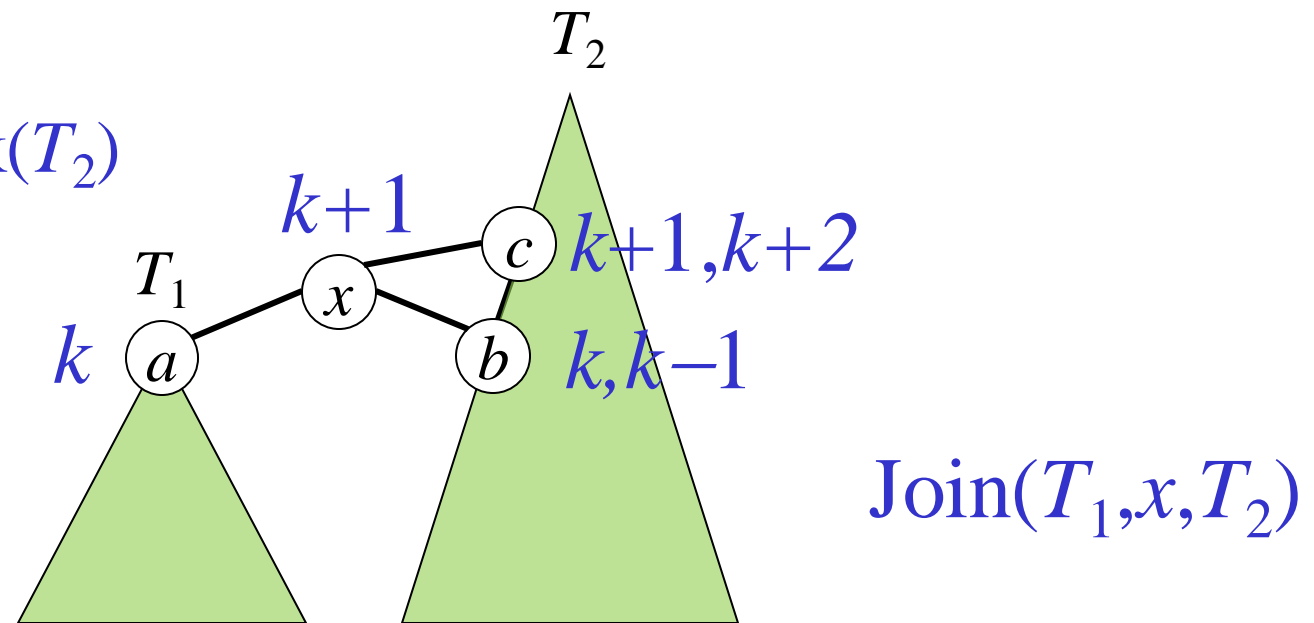
Suppose that all keys in T_1 are less than $x.key$
and that all keys in T_2 are greater than $x.key$

The tree formed is a valid search tree, but may be
very unbalanced, even if T_1 and T_2 are balanced

Joining two (W)AVL trees efficiently

Assume

$$\text{rank}(T_1) \leq \text{rank}(T_2)$$



b – first vertex on the left spine of T_2 with $\text{rank} \leq k$

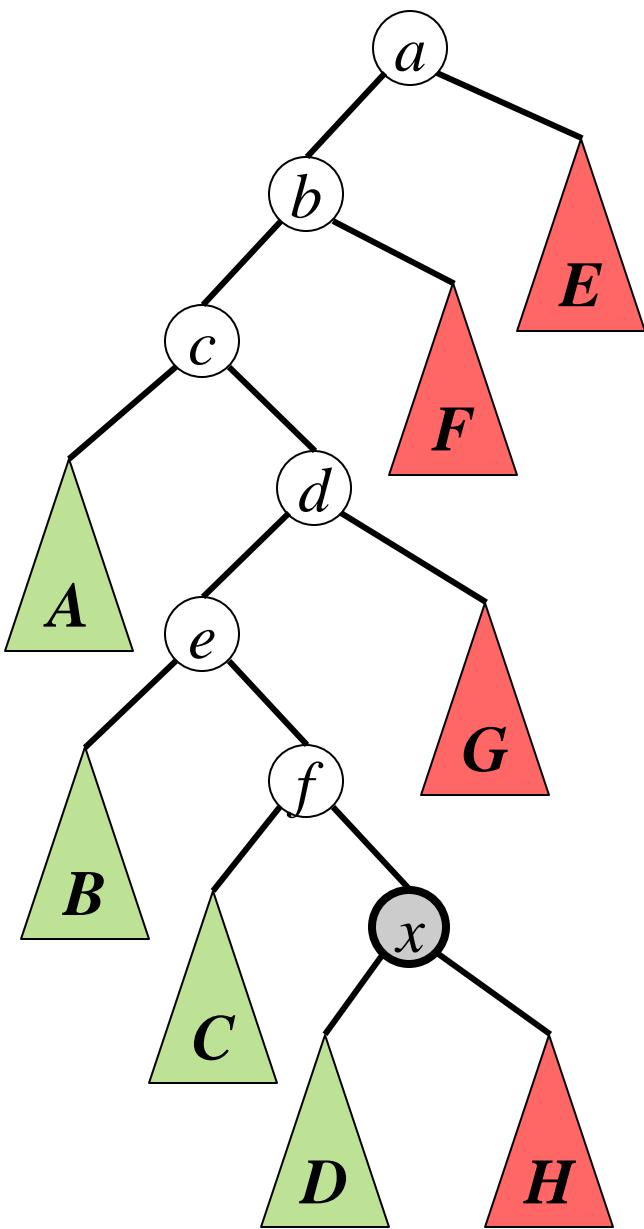
Do rebalancing from x , if needed

$O(\log n)$ time

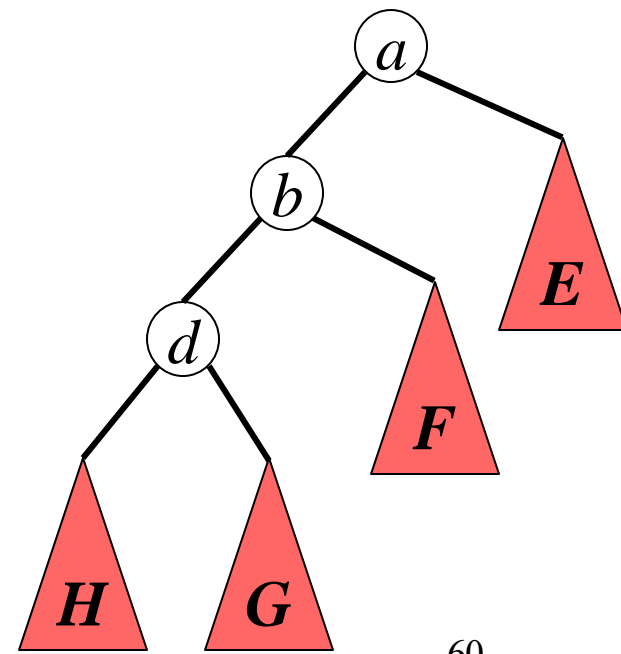
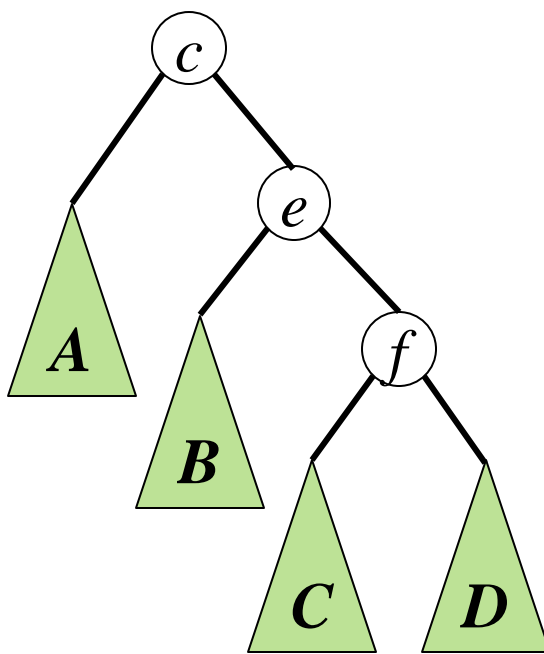
$O(\text{rank}(T_2) - \text{rank}(T_1) + 1)$ time

(if *ranks* maintained explicitly)

Splitting



x



Splitting with efficient joins

Suppose we need to join T_1, T_2, \dots, T_k
where $\text{rank}(T_1) \leq \text{rank}(T_2) \leq \dots \leq \text{rank}(T_k)$

Suppose $\text{rank}(\text{Join}(T_1, \dots, T_i)) \leq \text{rank}(T_i) + c$

$$\begin{aligned} & O \left(\sum_{i=2}^k \left| \text{rank}(T_i) - \text{rank}(\text{Join}(T_1, \dots, T_{i-1})) \right| + 1 \right) \\ &= O \left(\sum_{i=2}^k \text{rank}(T_i) - \text{rank}(T_1) + 1 \right) \\ &= O(\text{rank}(T_k) - \text{rank}(T_1) + k) = O(\log n) \end{aligned}$$

Rank and Select

Additional dictionary operations

$\text{Select}(D, i)$ – Return the i -th largest item in D
(indices start from 0)

$\text{Rank}(D, x)$ – Return the **RANK** of x in D ,
(i.e., the number of items x is larger than)

Can we still use **(W)AVL** trees?

Keep **sub-tree sizes**!

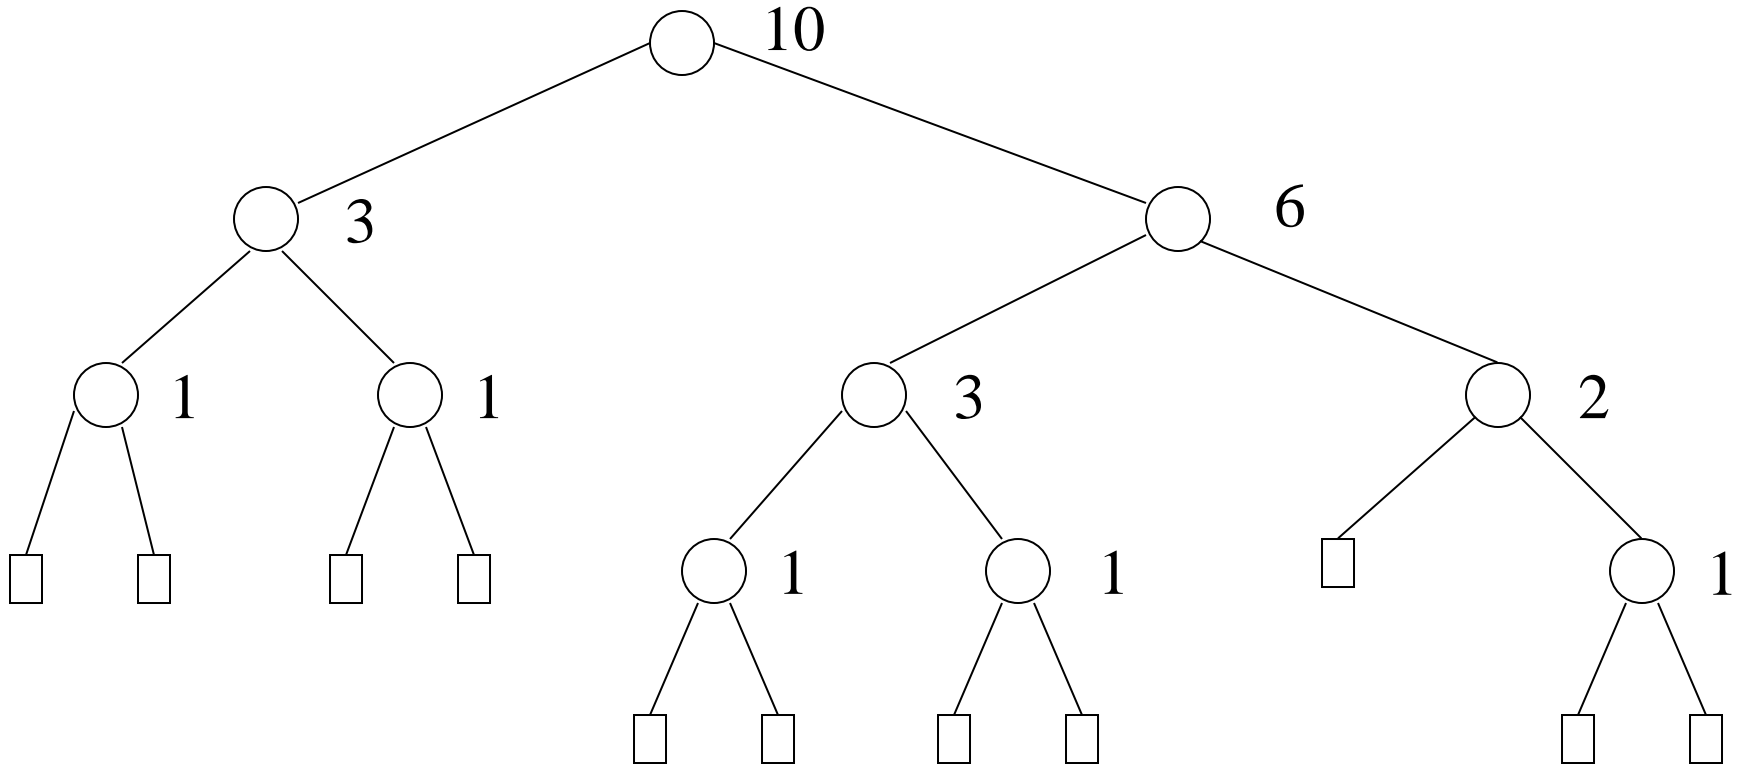
Caution!

Rank now has two meanings...

Which is unfortunate...

This is the established terminology...

Sub-tree sizes



$$x.size = x.left.size + x.right.size + 1$$

$$EXT.size = 0$$

Selection

Function $\text{Select}(x, i)$

$r \leftarrow x.\text{left.size}$

if $i = r$ **then**

 | **return** x

else if $i < r$ **then**

 | **return** $\text{Select}(x.\text{left}, i)$

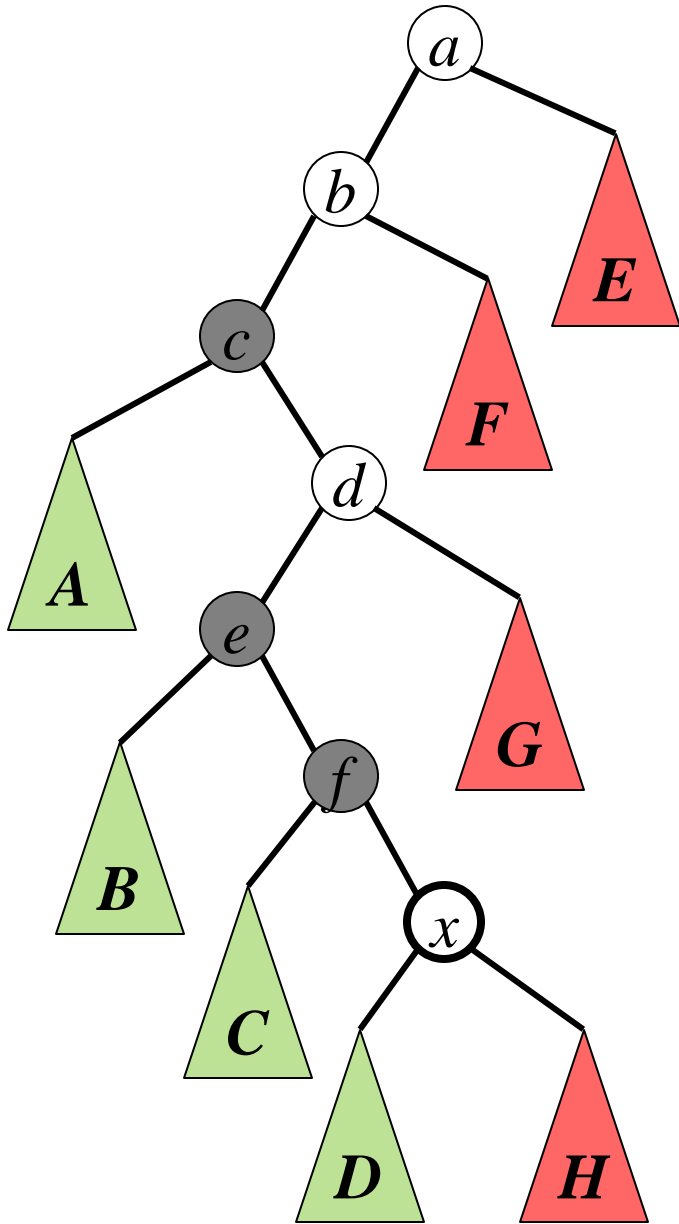
else

 | **return**

 | $\text{Select}(x.\text{right}, i - r - 1)$

Note: $0 \leq i < n$

RANK



$$\begin{aligned} \text{RANK}(x) = & \\ & (\text{size}(A) + 1) + \\ & (\text{size}(B) + 1) + \\ & (\text{size}(C) + 1) + \\ & \text{size}(D) \end{aligned}$$

RANK

Function Rank(T, x)

$r \leftarrow x.left.size$

$y \leftarrow x$

while $y \neq T.root$ **do**

if $y = y.parent.right$ **then**

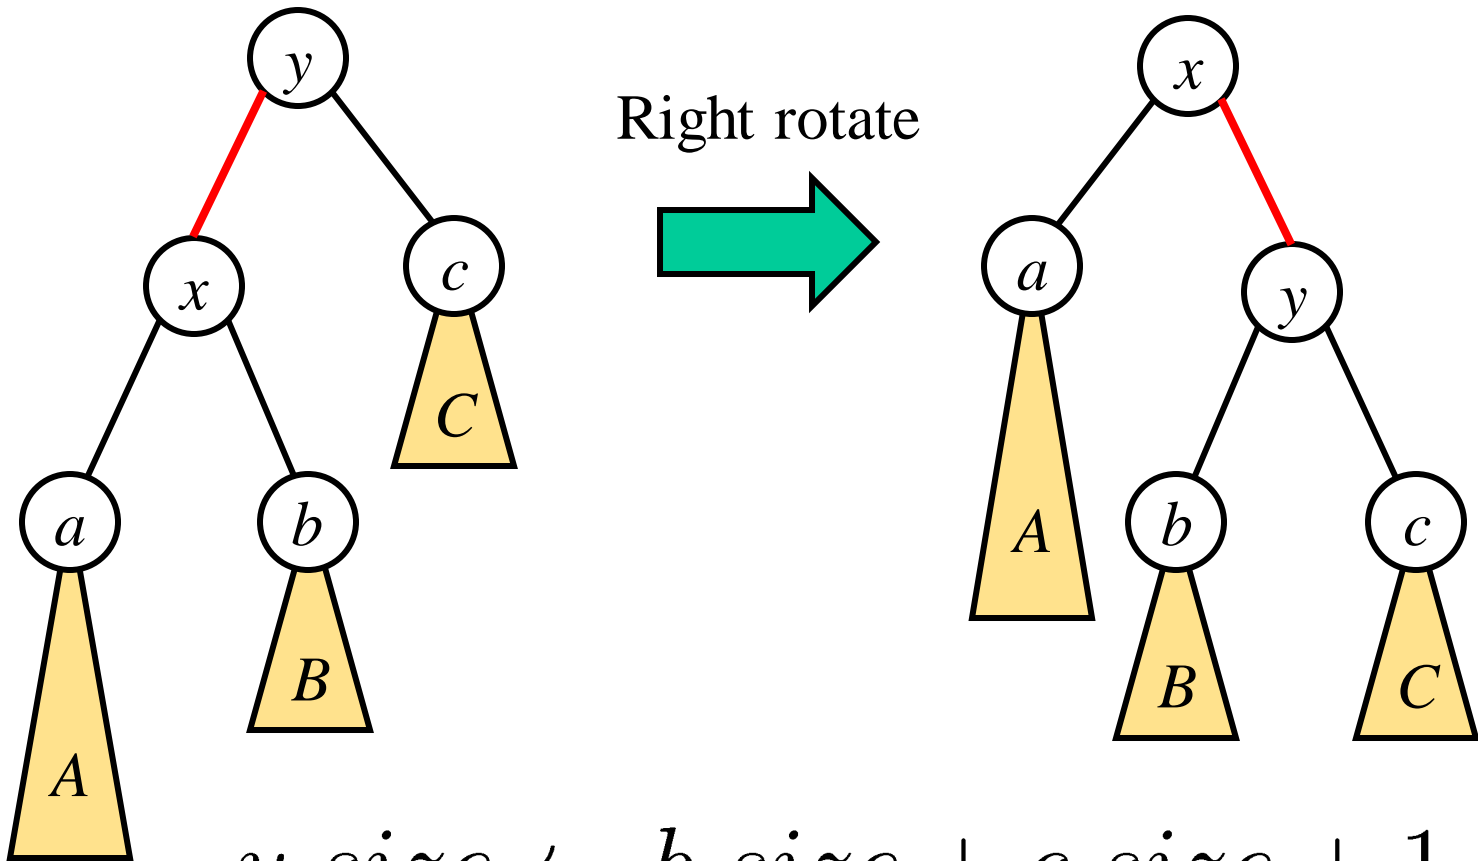
$r \leftarrow r + y.parent.left.size + 1$

$y \leftarrow y.parent$

return r

Recall that $EXT.size=0$

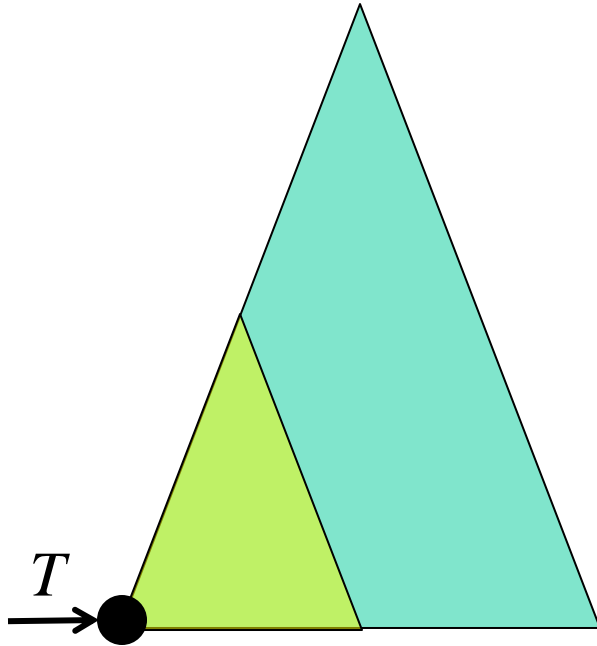
Easy to maintain **sizes**



$$y.size \leftarrow b.size + c.size + 1$$

$$x.size \leftarrow a.size + y.size + 1$$

Finger Search trees

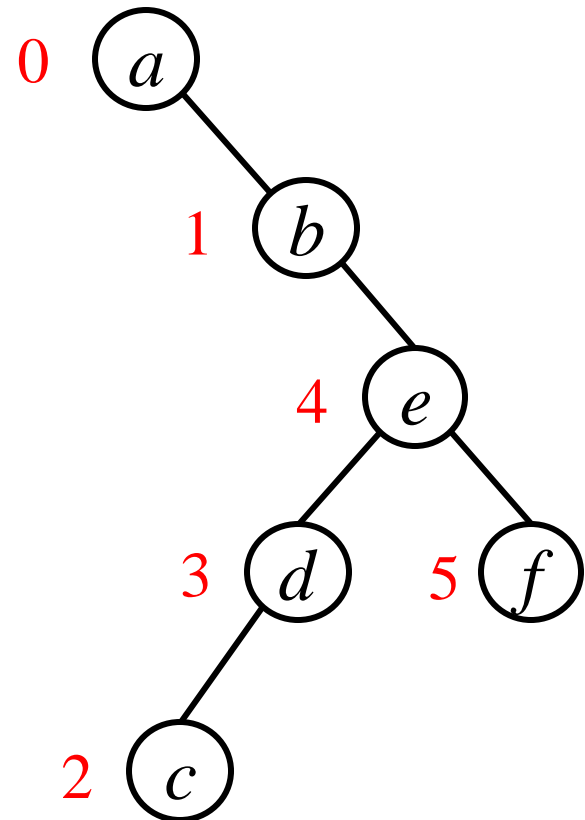
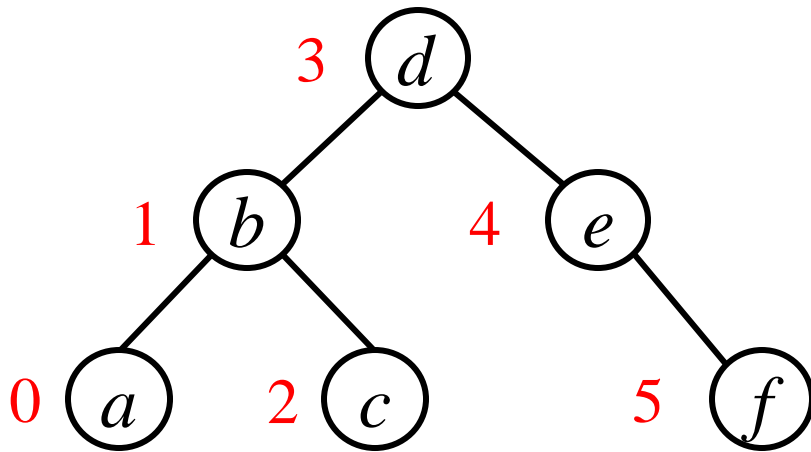


Maintain a pointer to the **minimum** element

Select(T, k) in
 $O(\log k)$ time

Lists as Trees

$[a\ b\ c\ d\ e\ f]$



Lists as Trees

Maintain the items in a tree:

i -th item in node of rank i

List-Node \rightarrow Tree-Node

Tree-Nodes have no explicit keys

(Implicitly maintained ranks play the role of keys)

Retrieve(i) \rightarrow Select(i)

Select, Insert-Rebalance and Delete-Rebalance
do not use keys

Lists as Trees: Insert

To insert a node z in the i -th position, where $0 \leq i < n$:

Find the current node of rank i .

If it has no **left** child, make z its **left** child.

Otherwise, find its predecessor
and make z its **right** child.

To insert a node z in the last position ($i = n$):

Find the last node and
make z its **right** child

Fix the tree

Lists as Trees: Delete

Delete a node z in the same way
a node is removed from a search tree

Implementation of lists

	Circular arrays	Doubly Linked lists	Balanced Trees
Insert/Delete-First/Last	$O(1)$	$O(1)$	$O(\log n)$
Insert/Delete(i)	$O(i+1)$	$O(i+1)$	$O(\log n)$
Retrieve(i)	$O(1)$	$O(i+1)$	$O(\log n)$
Concat	$O(n+1)$	$O(1)$	$O(\log n)$
Split(i)	$O(i+1)$	$O(i+1)$	$O(\log n)$

$O(i+1)$ can be replaced by $O(\min\{i+1, n-i\})$

Splay Trees (Self-adjusting trees)

[Sleator-Tarjan (1983)]

Do not maintain any balance!

When a node is accessed, **splay** it to the root

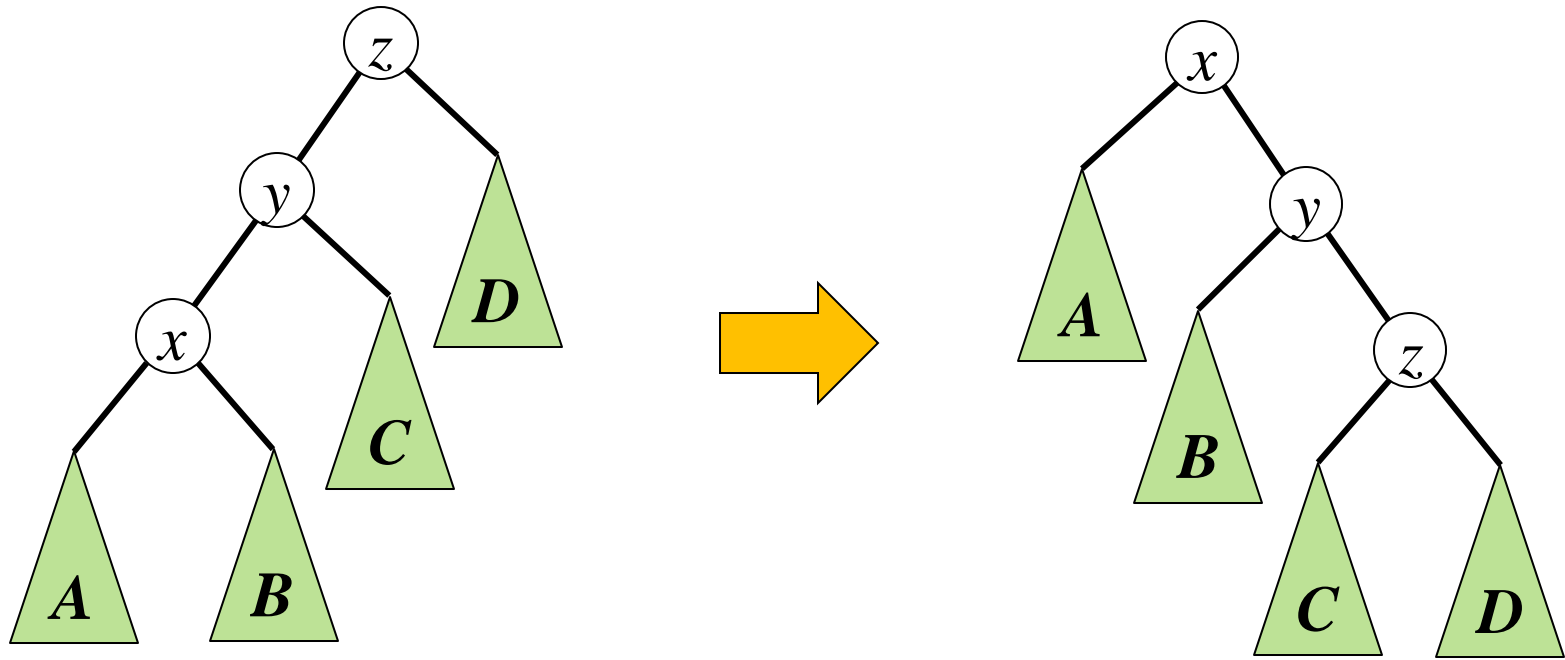
A node is splayed using a sequence of
zig-zig and **zig-zag** steps

Amortized cost of each operation is $O(\log n)$

Total cost of n operation is $O(n \log n)$

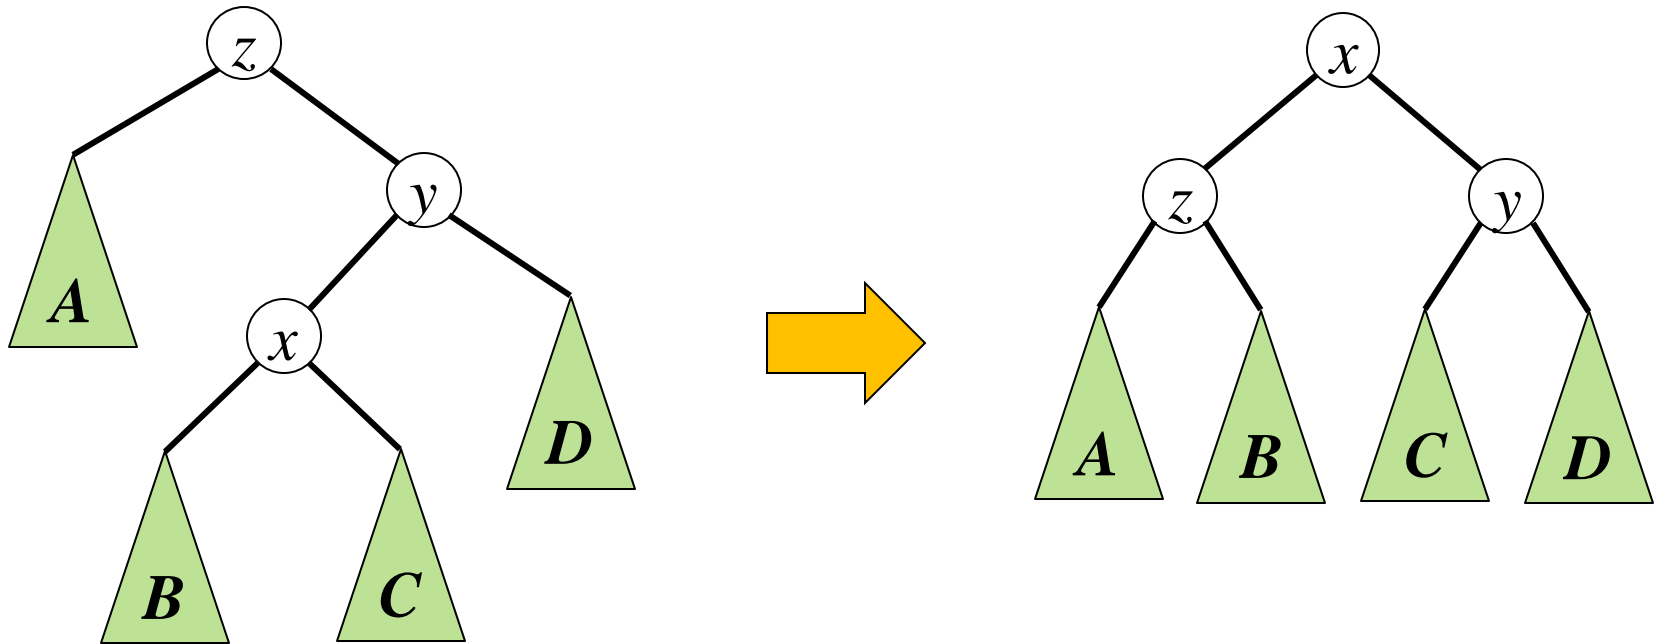
Many other amazing properties

Zig-Zig



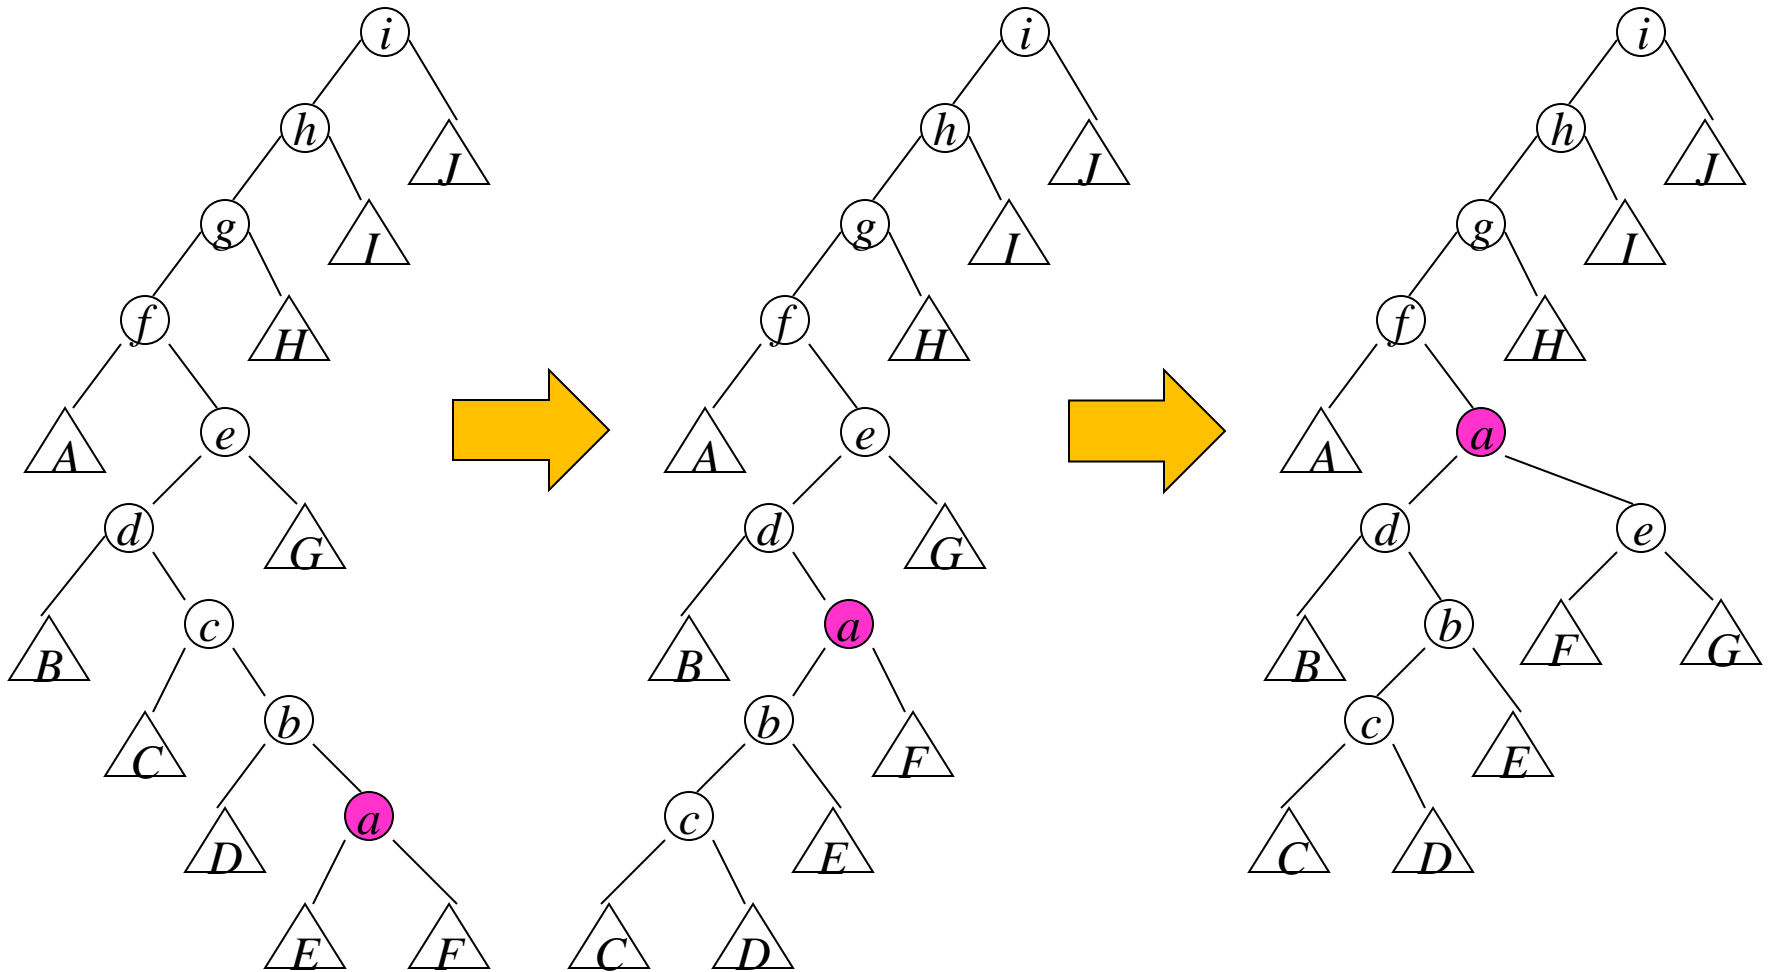
Rotate y - z left, then rotate x - y left

Zig-Zag

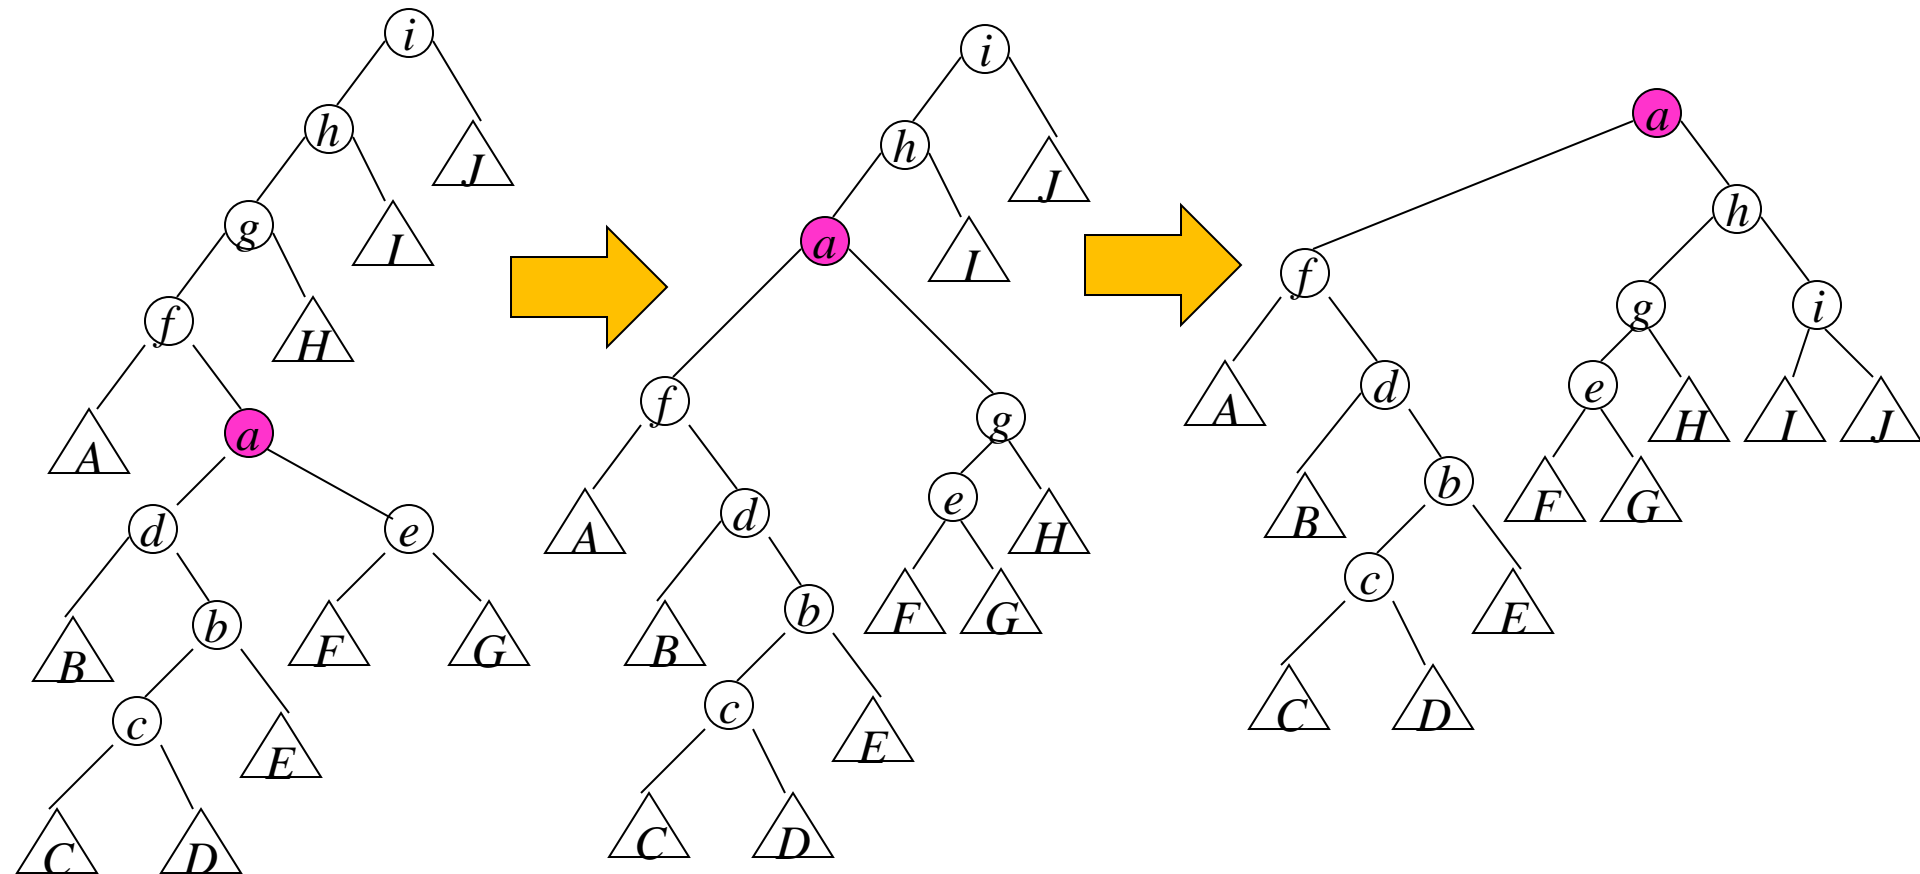


Rotate x - y right, then rotate x - z left

Splaying (example)



Splaying (example cont)



Splay Trees (Self-adjusting trees)

[Sleator-Tarjan (1983)]

Amortized cost of each operation is $O(\log n)$

Total cost of n operation is $O(n \log n)$

Many other amazing properties

Some intriguing open problems

Play with them yourself:

<http://webdiis.unizar.es/asignaturas/EDA/AVLTree/avltree.html>